

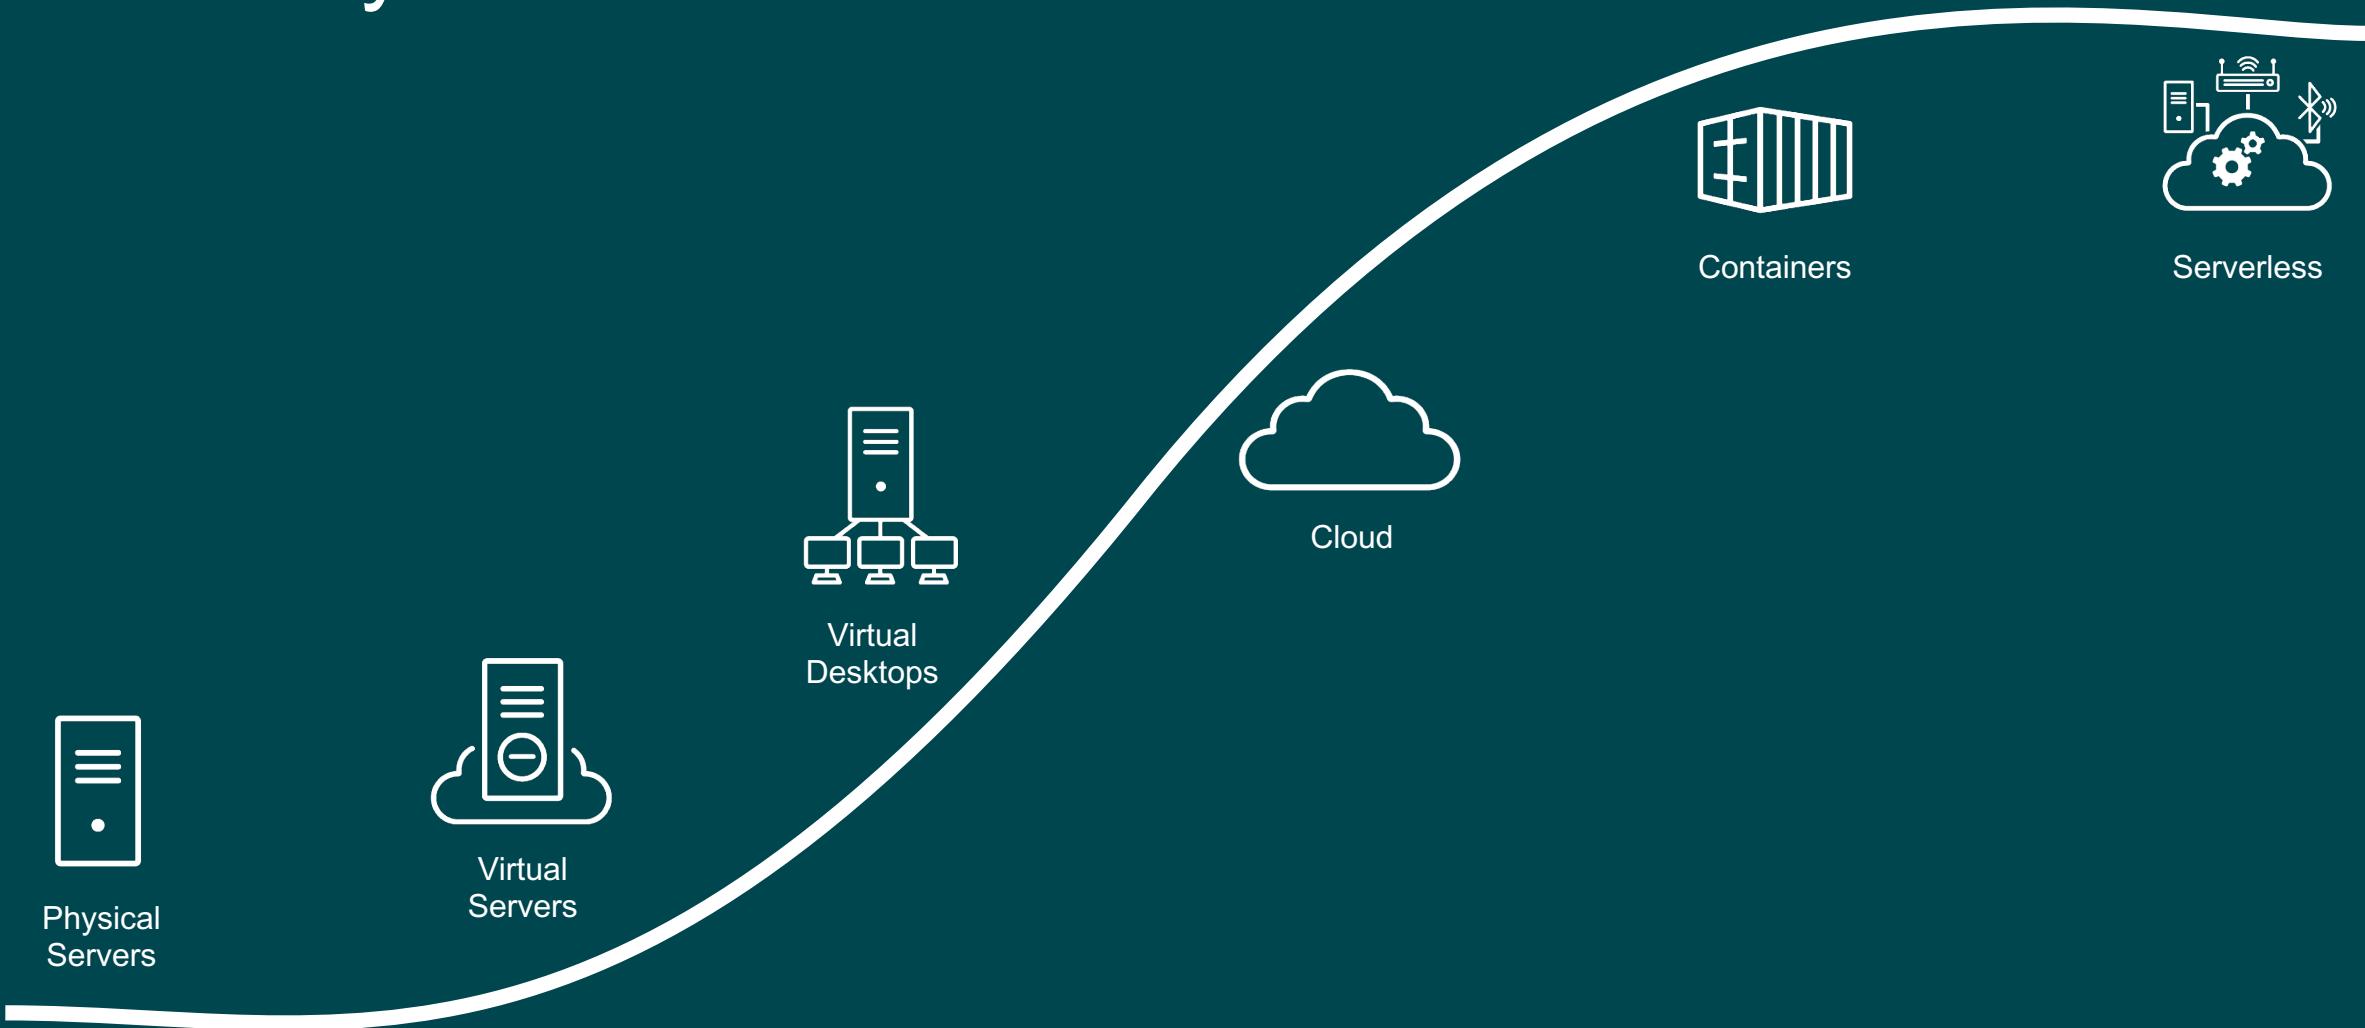
Container Security

Tejas Sheth
Sr. Security Solutions architect

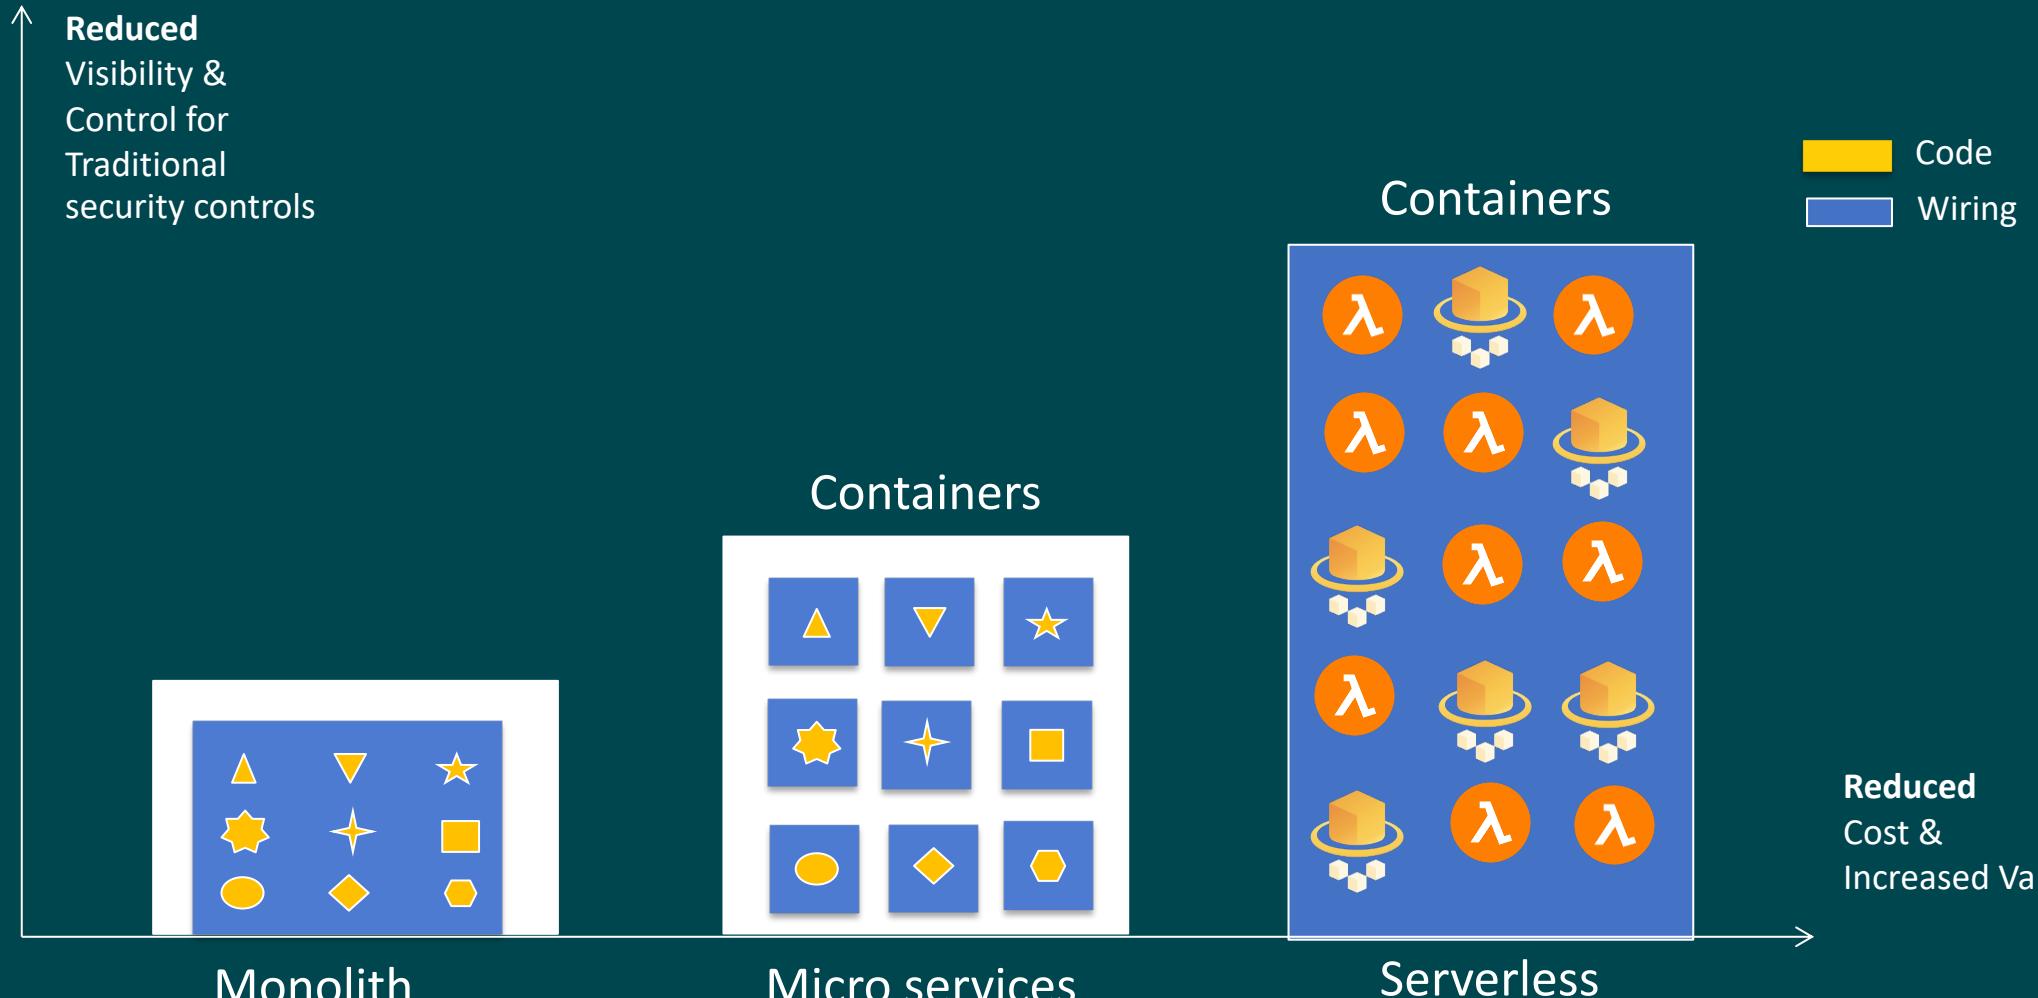
Objective

- Container security Threat vectors
- Container security controls
 - Preventive controls
 - Risk detection and remediation at early stage for SDLC
 - Risk remediation trade offs
 - Best practices
 - Protective controls
 - Threat mitigation at edge
 - Threat mitigation at VPC level
 - Threat mitigation at container level
 - Threat detection and response controls
- Dive deep on each control with container security best practice
- Security automation use cases

Journey to abstraction



Abstraction Tread-off



**“Gartner estimates that 90%
of global organizations will be
running containerized
applications in production by
2026”**

¹ Gartner, [The Innovation Leader’s Guide to Navigating the Cloud-Native Container Ecosystem](#), Arun Chandrasekaran, Wataru Katsurashima, August 18, 2021.

AWS managed container services

Management

Deployment, scheduling,
scaling, and management of
containerized applications



**Amazon Elastic
Container Service**



**Amazon Elastic
Kubernetes Service
(EKS)**

Hosting

Where the containers run



Amazon EC2



AWS Fargate

Image registry

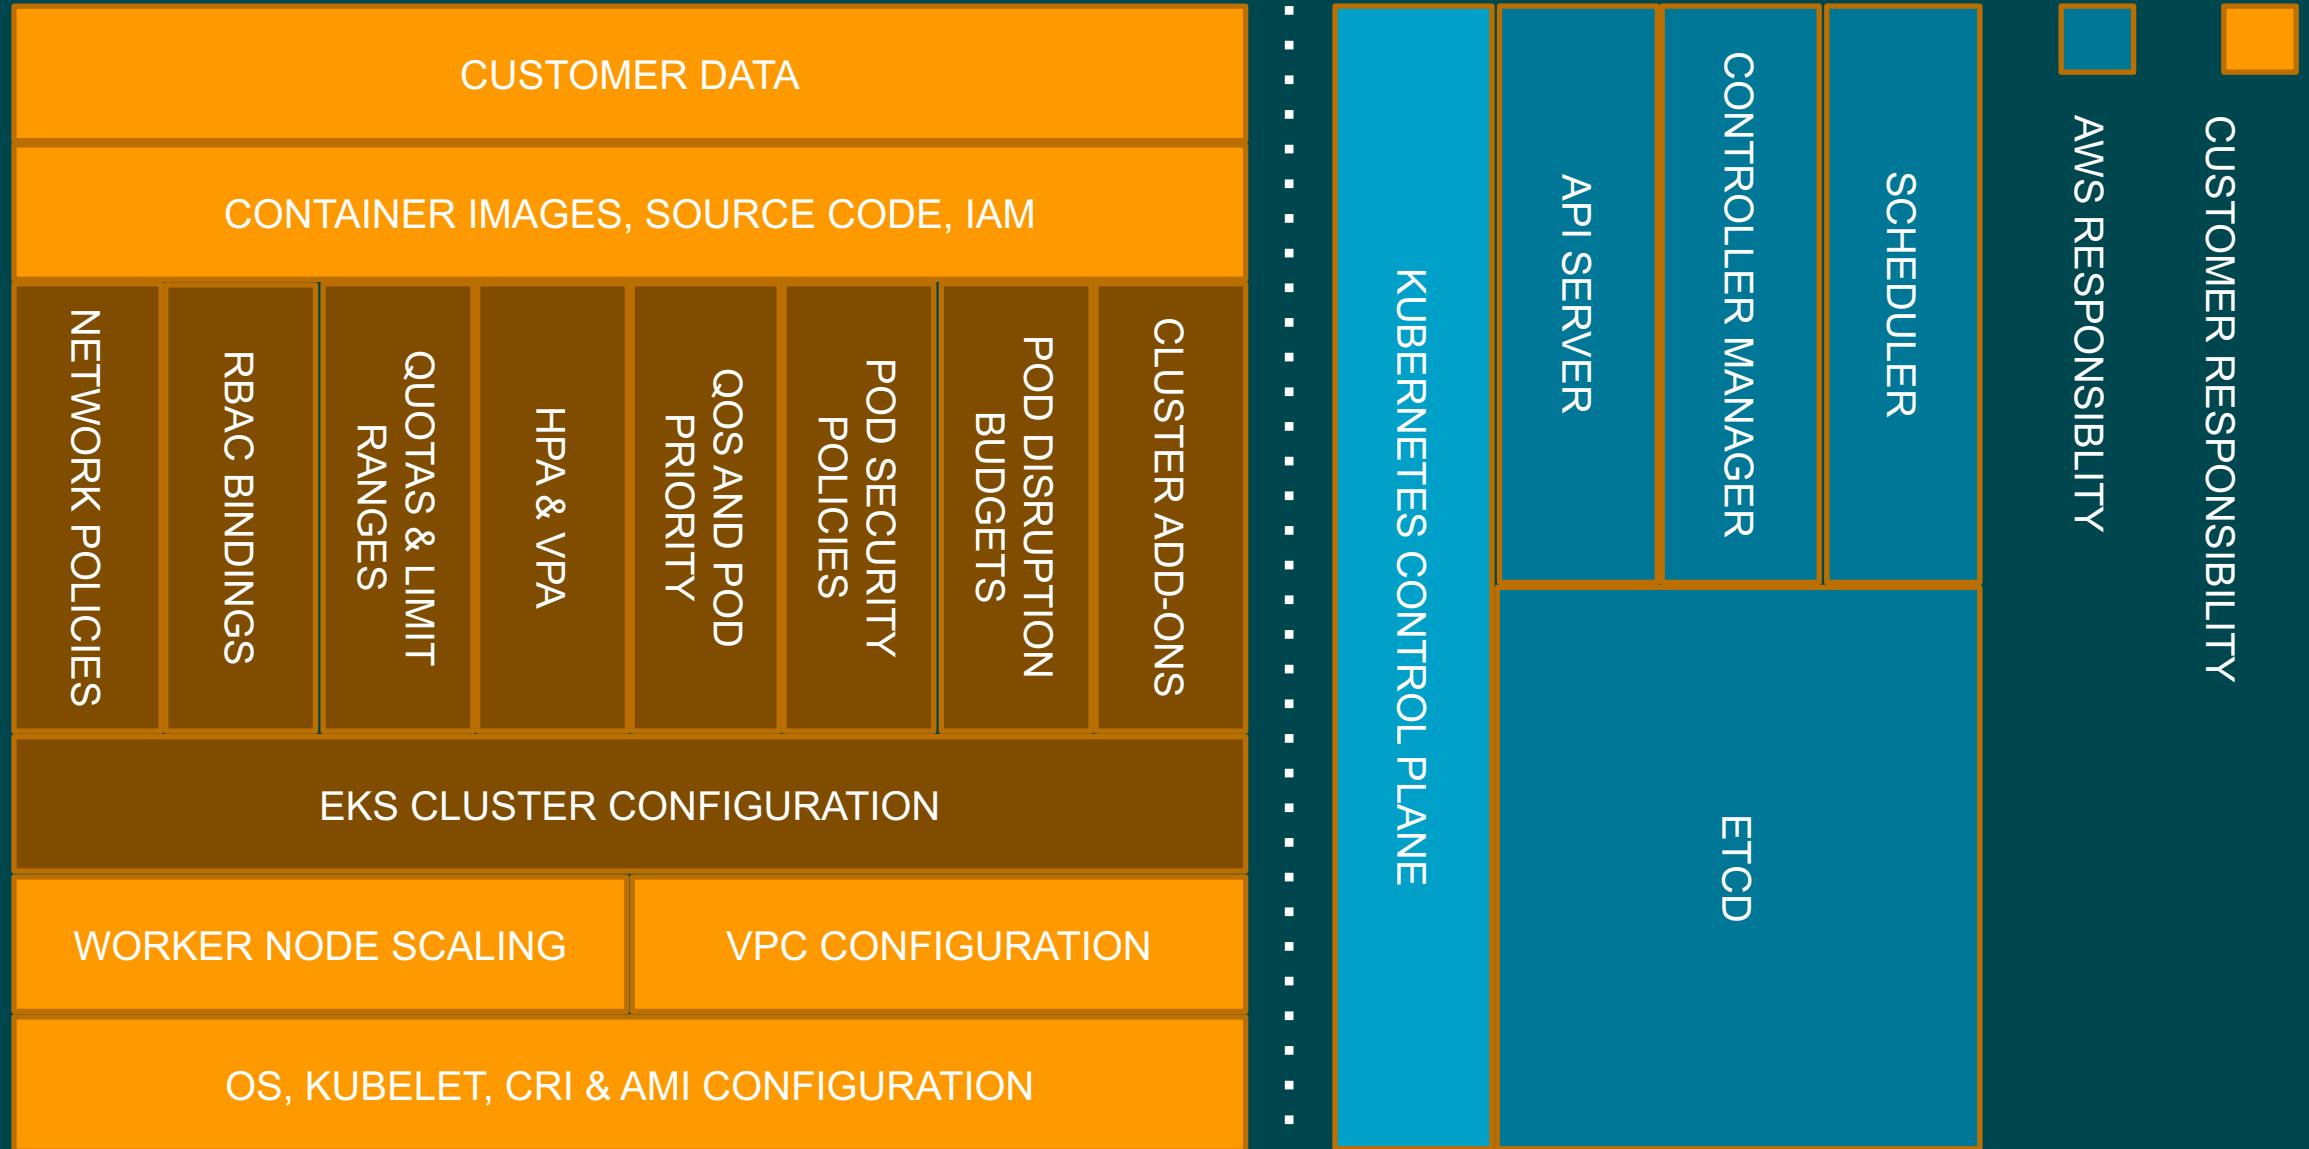
Container image repository



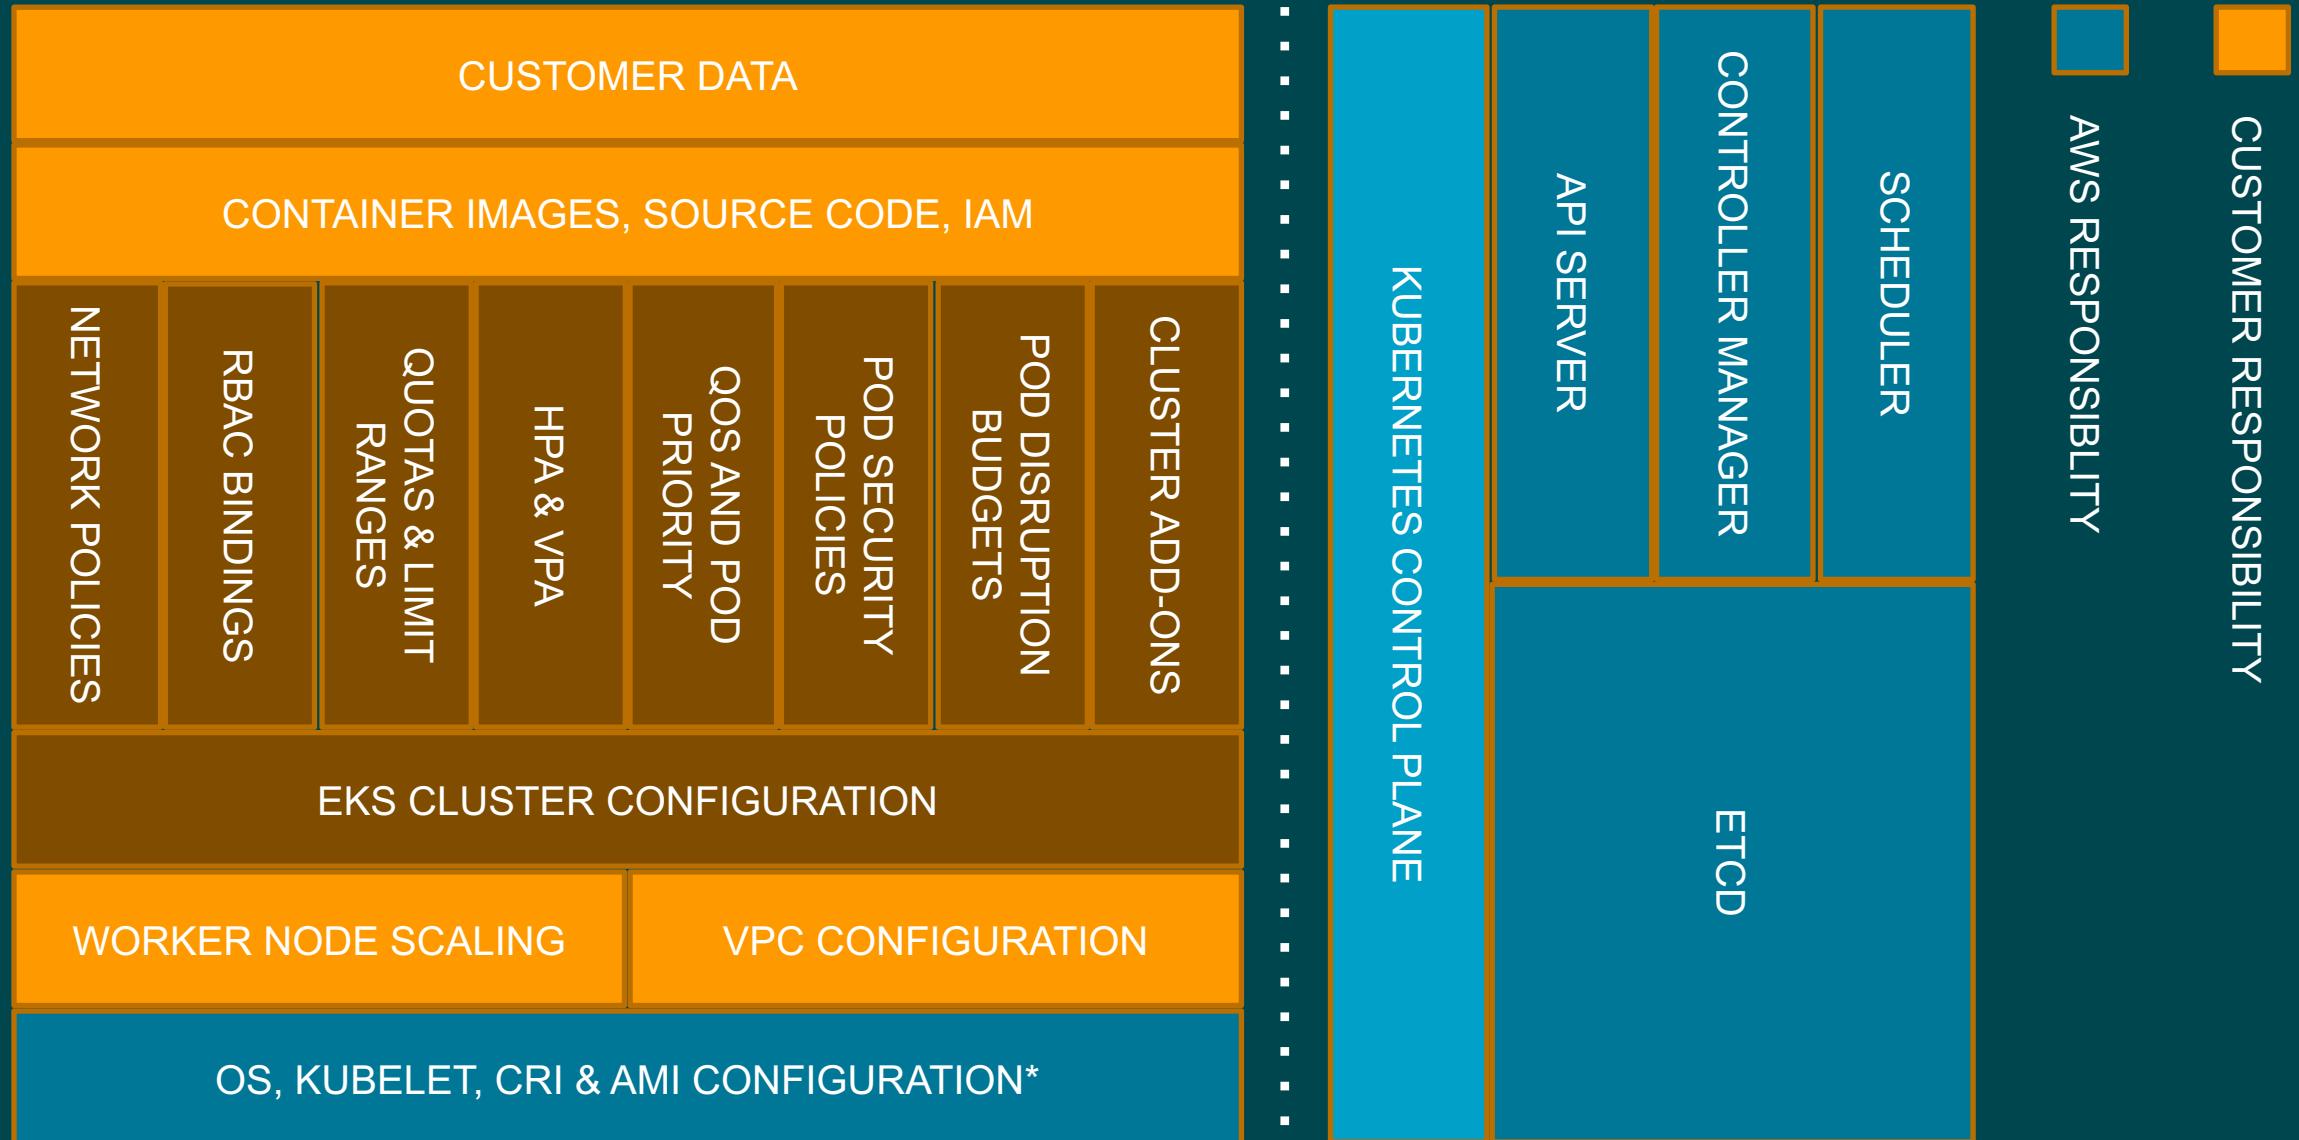
**Amazon Elastic
Container Registry**



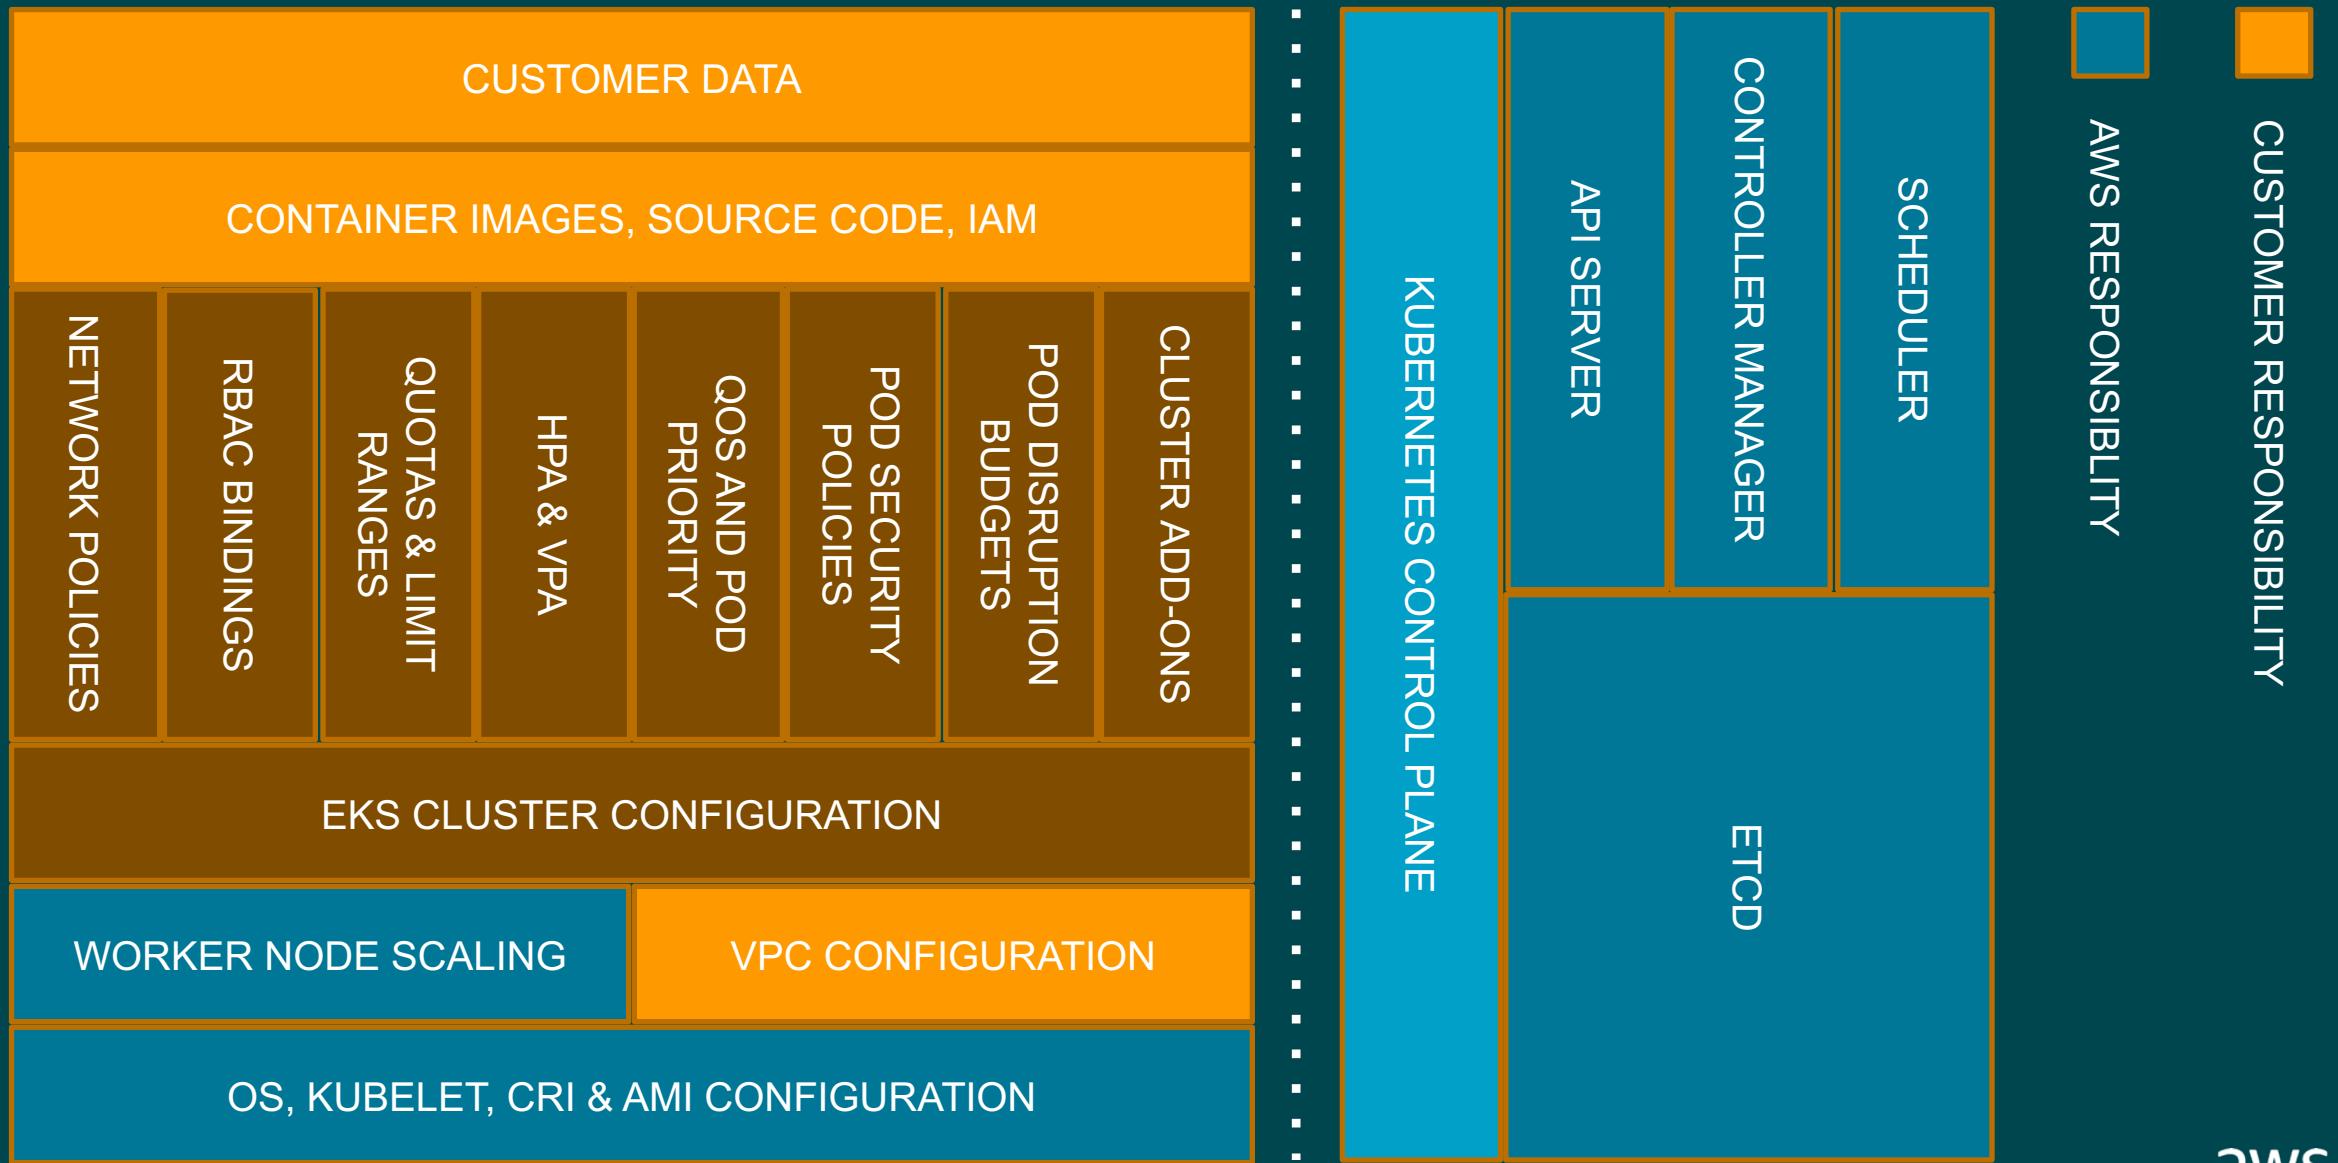
Amazon EKS with self-managed workers



Amazon EKS with managed node groups



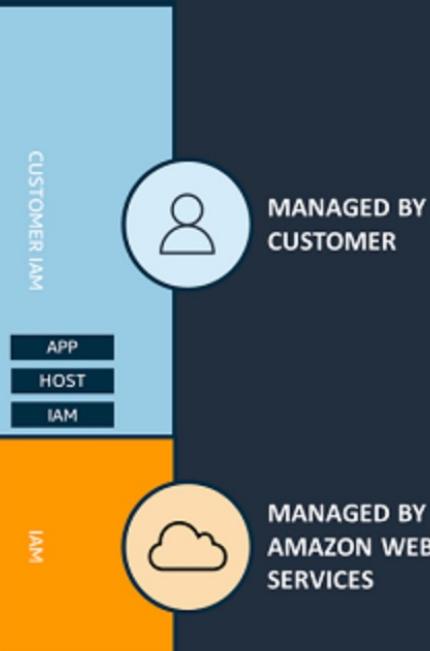
Amazon EKS on AWS Fargate



AWS ECS shared responsibility model

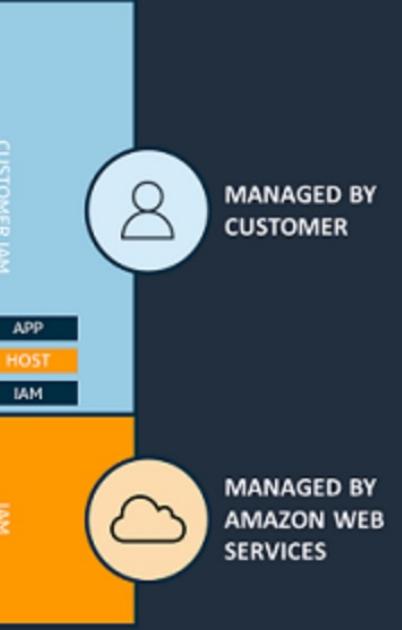
Amazon ECS on EC2

ECS AGENT	APPLICATION				
	CONTAINER	HARDENING	MONITORING	PATCHING	
TASK					
WORKER NODE CONFIGURATION	HARDENING	MONITORING	PATCHING		
NETWORK CONFIGURATION	NACLS	SECURITY GROUPS	ROUTE TABLES	VPC	
DATA	CLIENT-SIDE ENCRYPTION	SERVER-SIDE ENCRYPTION	NETWORK TRAFFIC PROTECTION		
ECS CONTROL PLANE					
VPC ENDPOINTS	FOUNDATION SERVICES	COMPUTE	STORAGE	DATABASES	NETWORKING
	AWS GLOBAL INFRASTRUCTURE	REGIONS	ZONES	EDGE LOCATIONS	



Amazon ECS on Fargate

ECS AGENT	APPLICATION				
	CONTAINER	HARDENING	MONITORING	PATCHING	
TASK					
WORKER NODE CONFIGURATION	HARDENING	MONITORING	PATCHING		
NETWORK CONFIGURATION	NACLS	SECURITY GROUPS	ROUTE TABLES	VPC	
DATA	CLIENT-SIDE ENCRYPTION	SERVER-SIDE ENCRYPTION	NETWORK TRAFFIC PROTECTION		
ECS CONTROL PLANE					
VPC ENDPOINTS	FOUNDATION SERVICES	COMPUTE	STORAGE	DATABASES	NETWORKING
	AWS GLOBAL INFRASTRUCTURE	REGIONS	ZONES	EDGE LOCATIONS	



Container risks

01 Vulnerable application code

02 Poorly configured containers and pods

03 Insecure networking

04 Malware

05 Unauthorized behavior

06 Secrets exposure

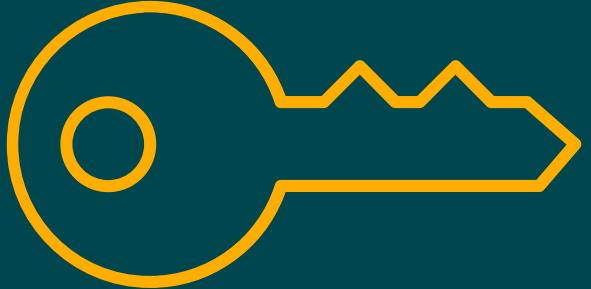




Are you working on a current AWS opportunity for Container security?

- Yes
- No

Security controls



Preventative



Detective

Risk mitigation techniques



Preventive
controls



Protective
controls



Threat
Detection and
Response



Preventive controls

Container security foundations

Preventive controls



Adopt a multi account architecture



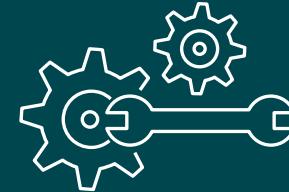
Establish a centralized logging account



Implement the principle of least privilege



Protect data in transit and at rest

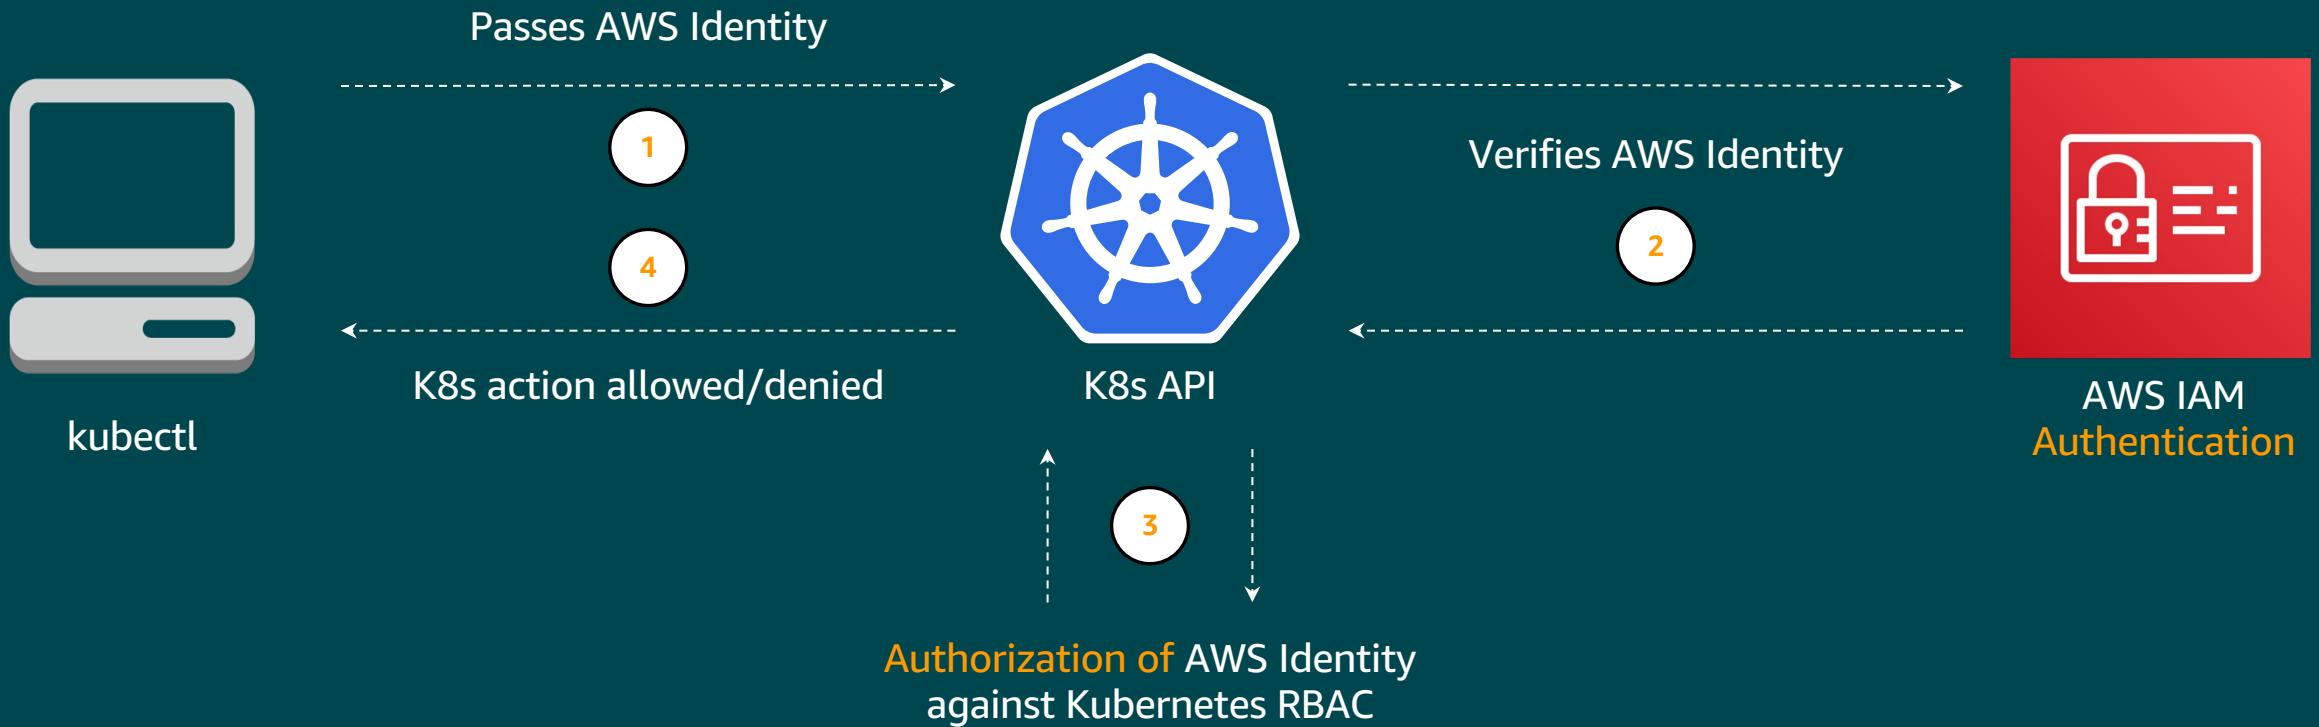


Automate tasks to save time and reduce risk



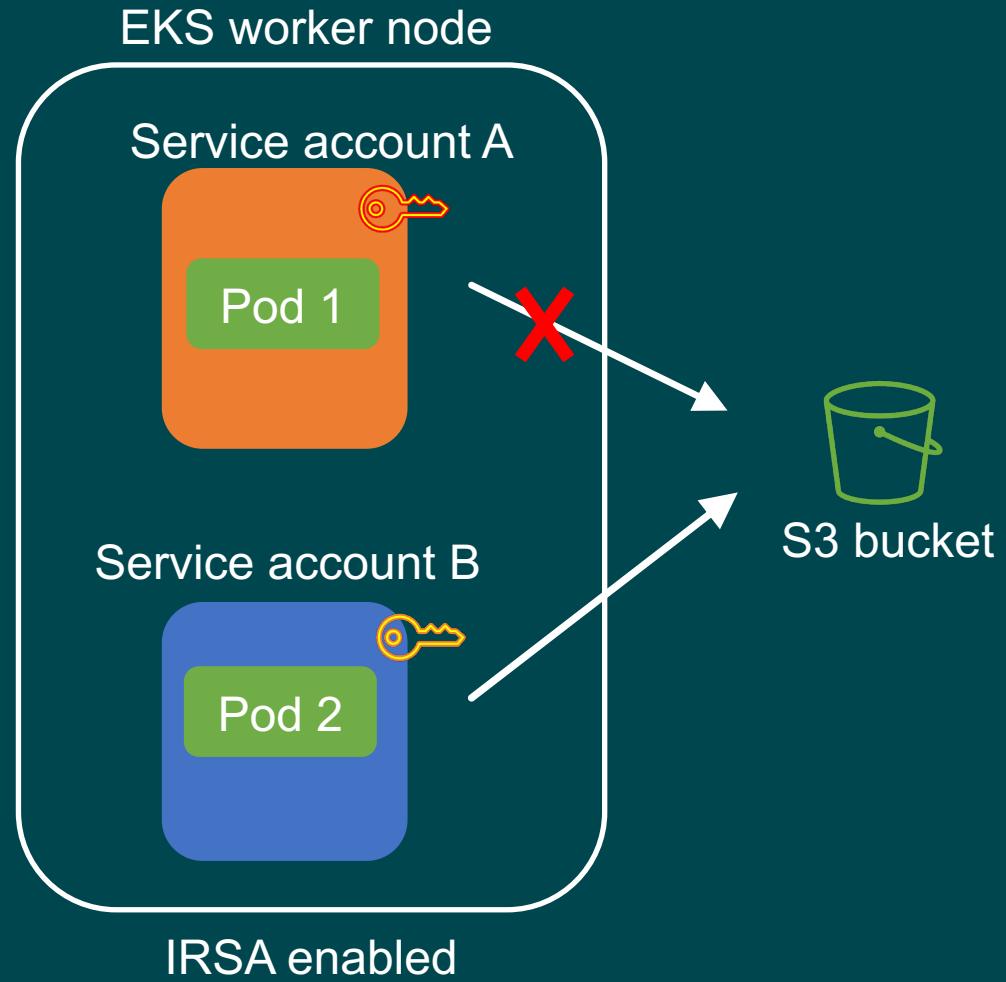
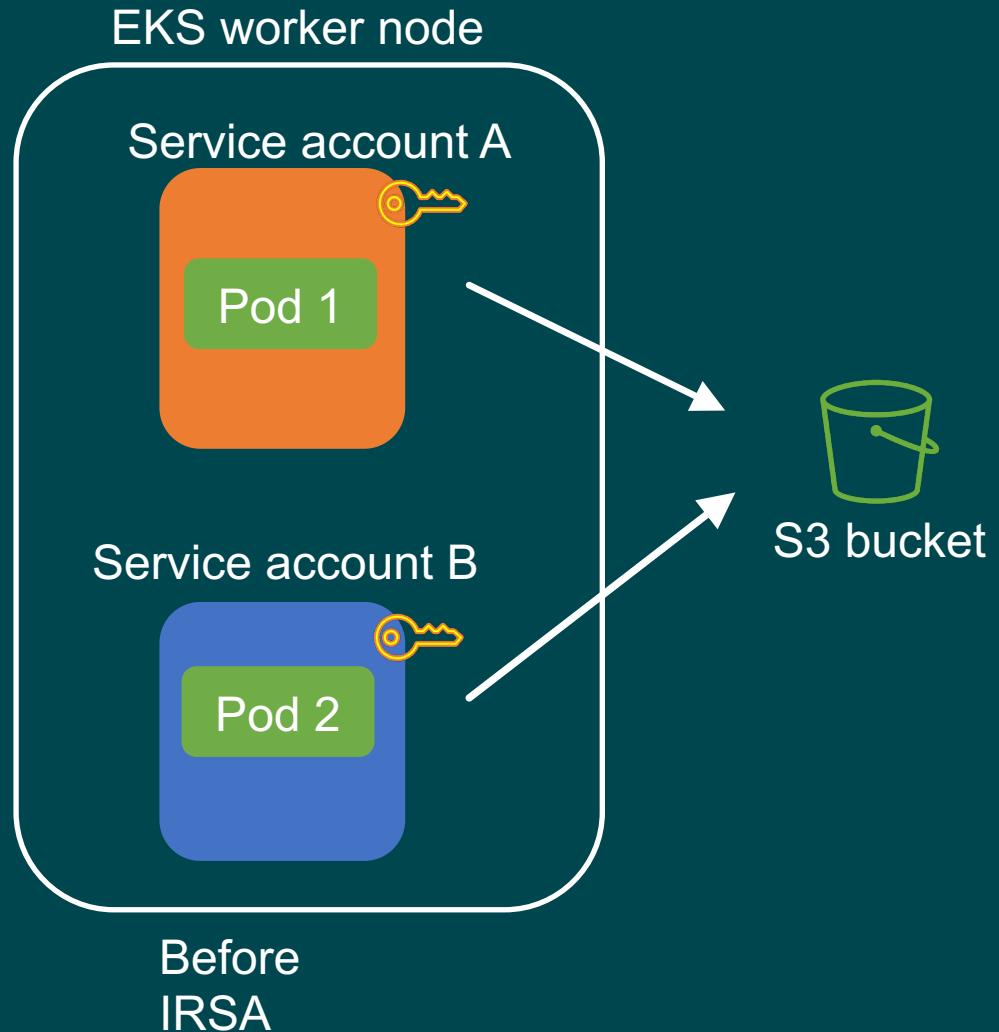
Apply security at all layers

Amazon EKS: IAM Authentication + kubectl



Container IAM – Best practices

ENABLE PERMISSION PER POD INSTEAD OF PER NODE



Network Security

VPC

- No changes to VPC (non-dedicated)
- Traffic routed through ENI to task

Security Group

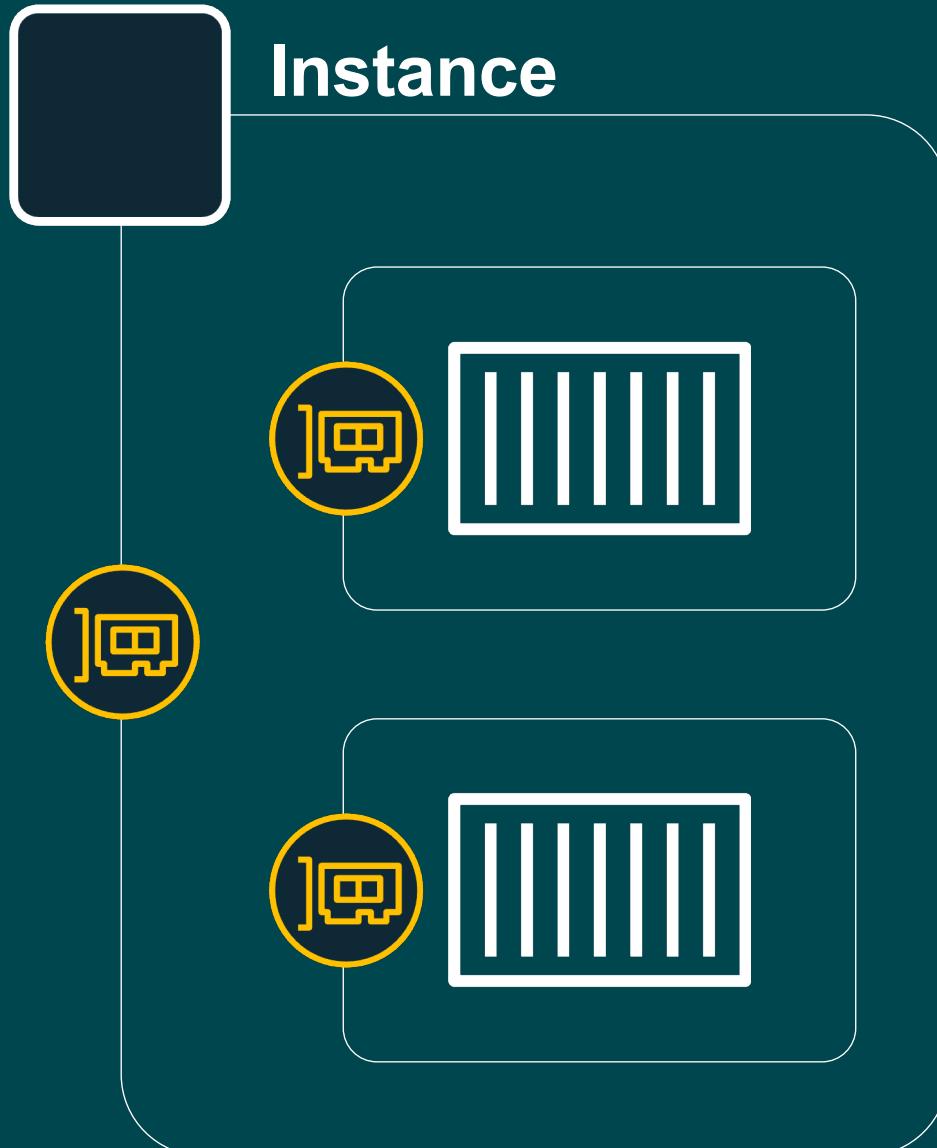
- Isolated security group per task
- Task-level specification

Task Network

- Uses CNI plugin architecture
- Isolated network namespace per task

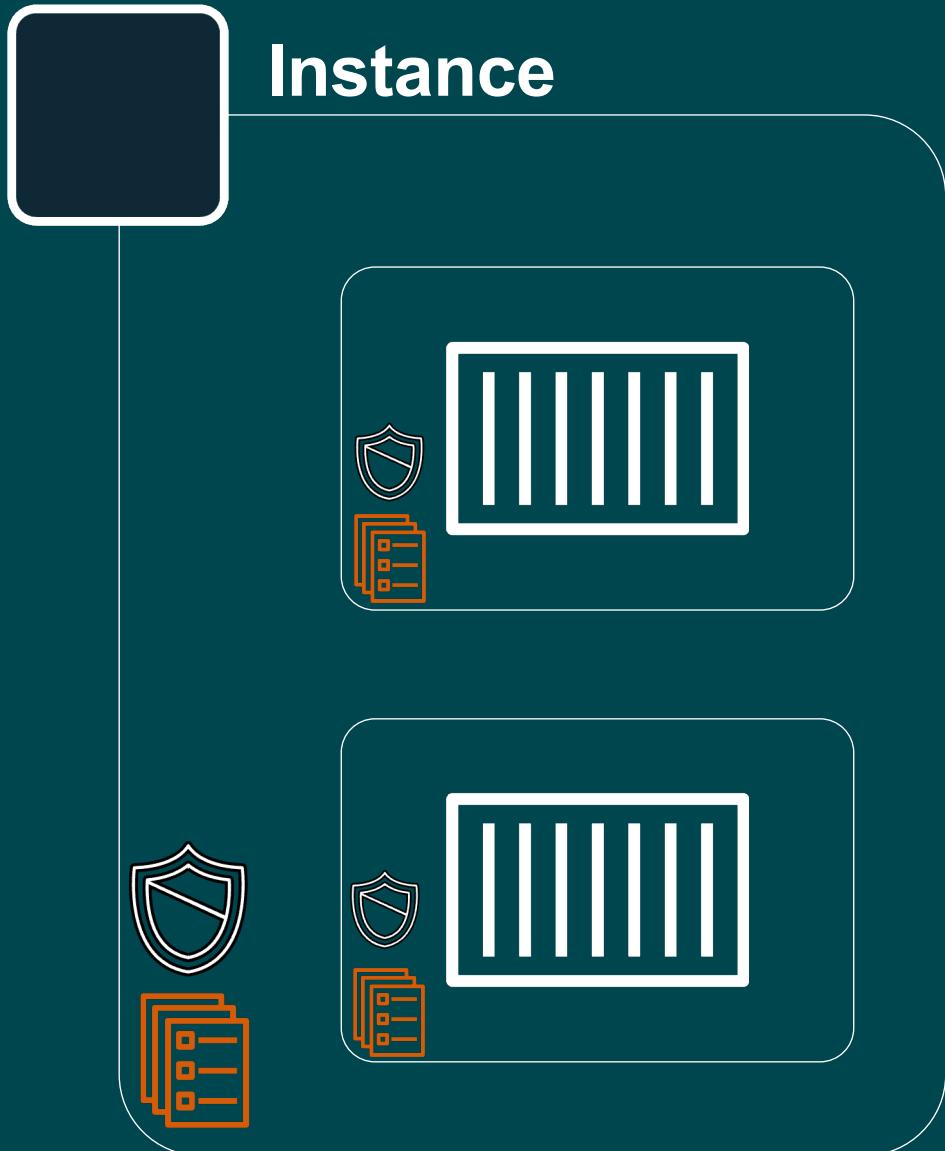
PrivateLink

- Pull images from ECR within VPC



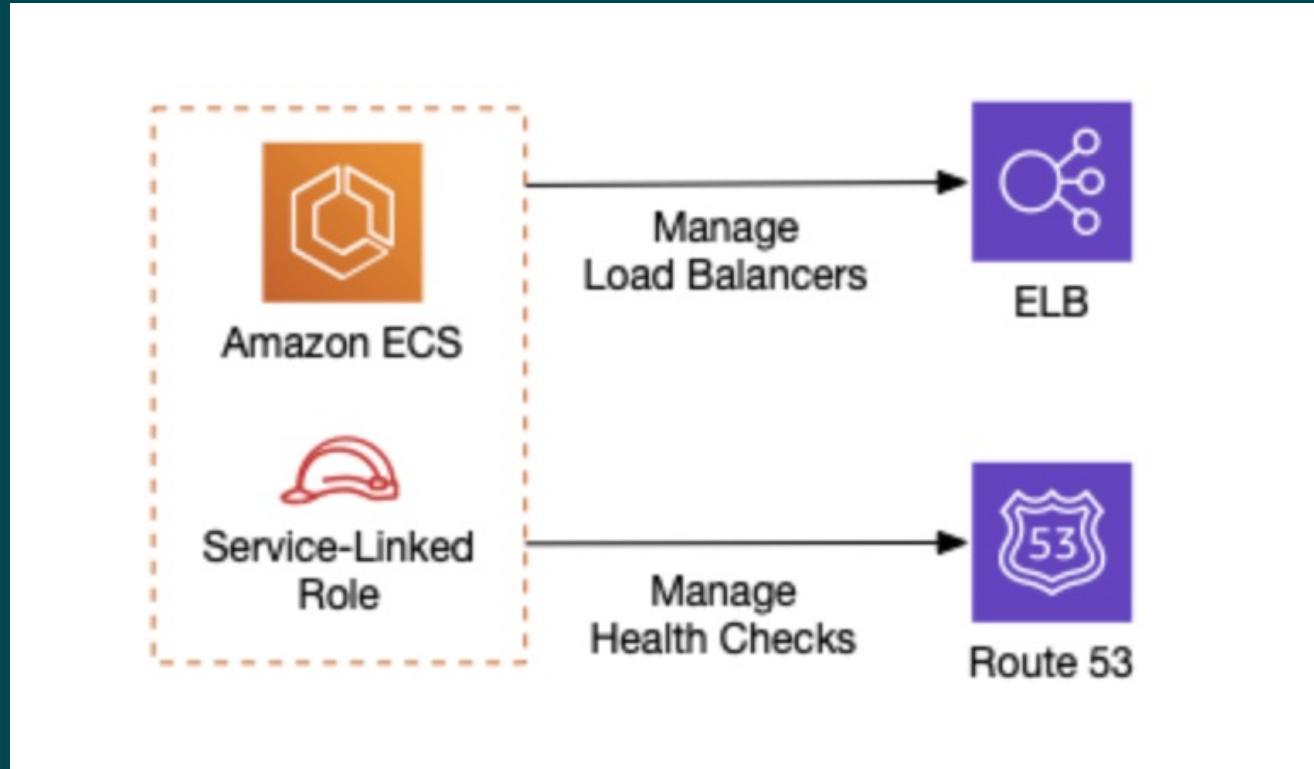
K8S Security hardening

- Define clusters based on workloads
- Kubernetes Node Affinity
- Kubernetes Pod limits quotas
- Kubernetes Network policies
- Node hardening SELinux
- Secure computing (Seccomp) to prevent unwanted system calls from container



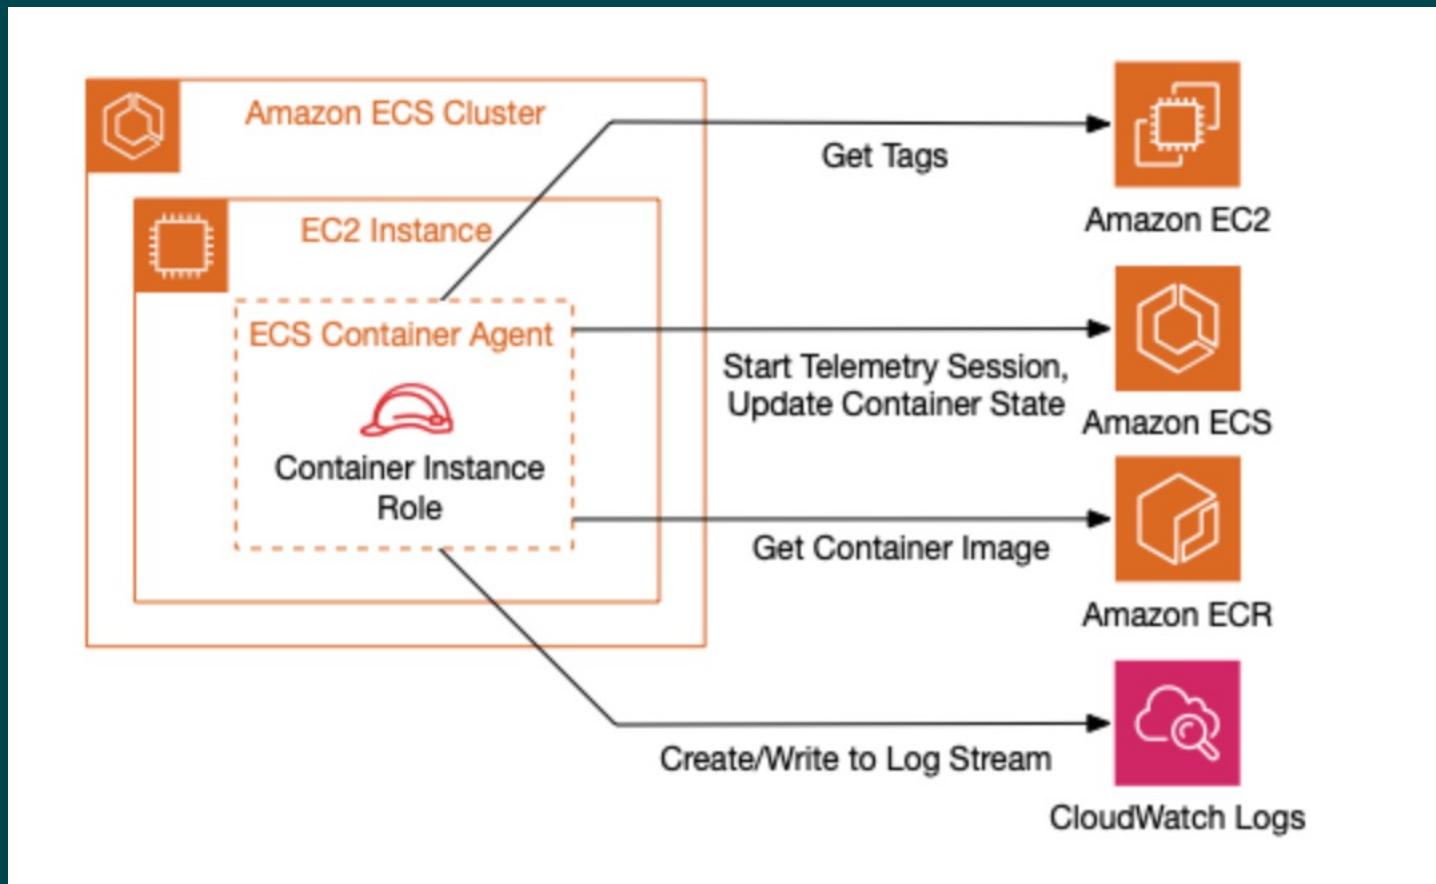
Service-Linked Role

Preventive Controls



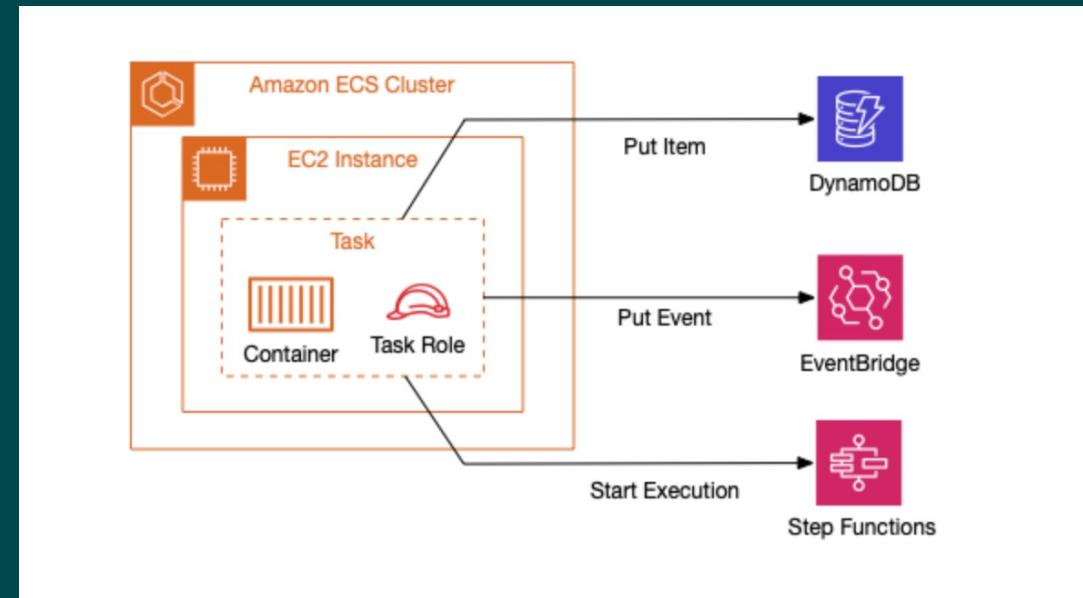
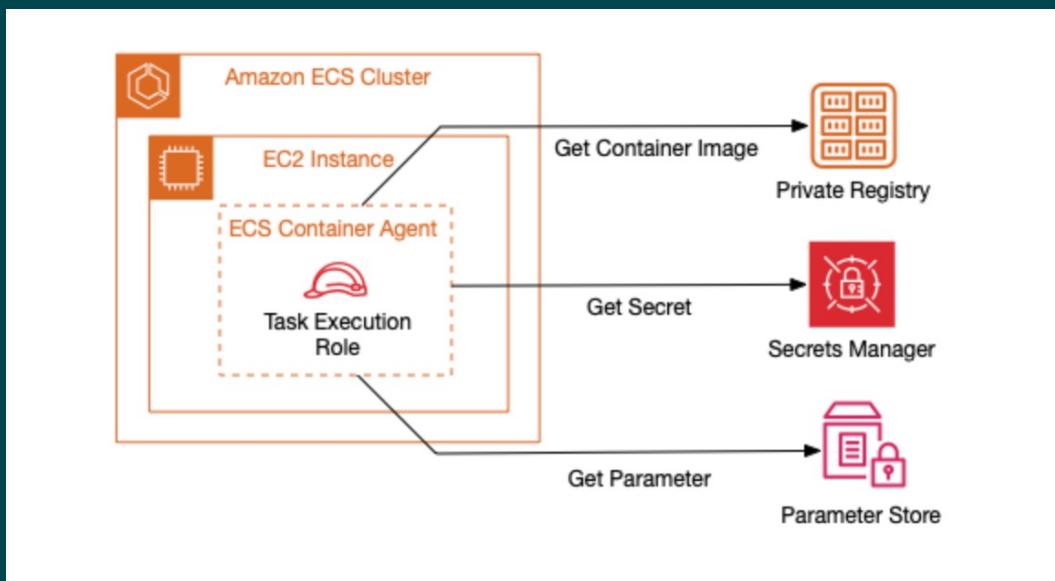
Container Instance Role

Preventive Controls



Task Execution Role

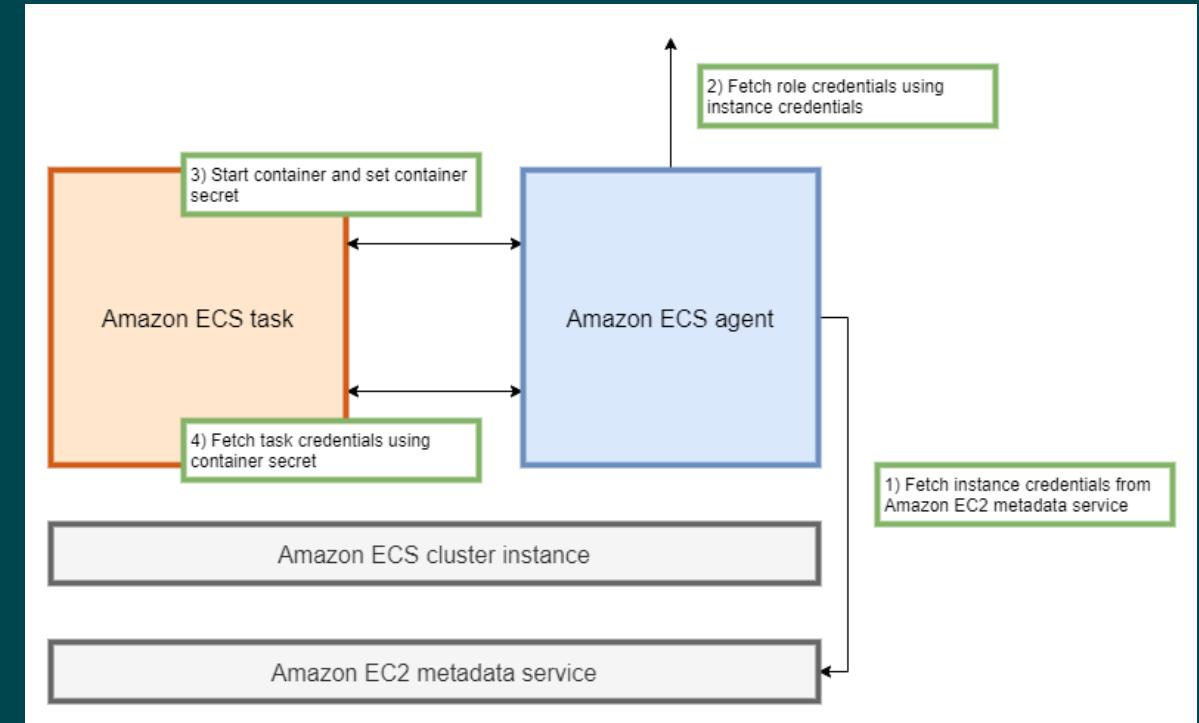
Preventive Controls



ECS Metadata services access with task role

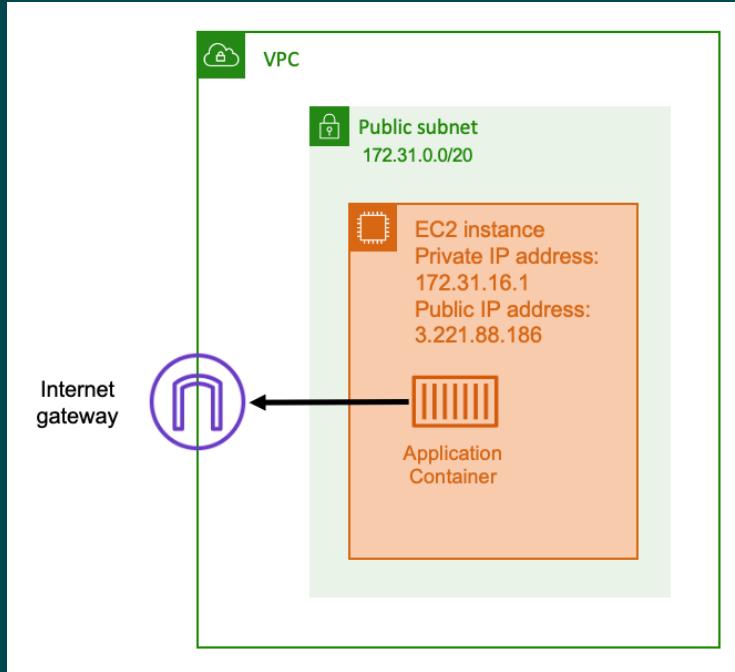
Preventive Controls

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Sid": "",  
      "Effect": "Allow",  
      "Principal": {  
        "Service": "ecs-tasks.amazonaws.com"  
      },  
      "Action": "sts:AssumeRole"  
    }  
  ]  
}
```



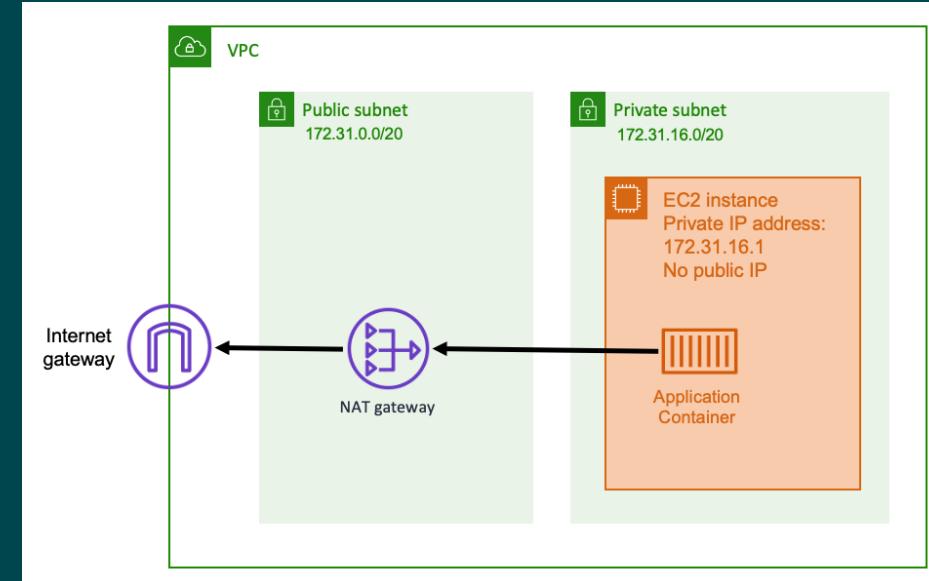
Connecting to the internet

Preventive Controls



Public Subnet

- Pay CLOSE ATTENTION to your security group and firewall rules. Make sure SSH port 22 is not open!
- Use case: public application with requirements of large amounts of bandwidth or minimal latency
- E.g., video streaming, gaming services

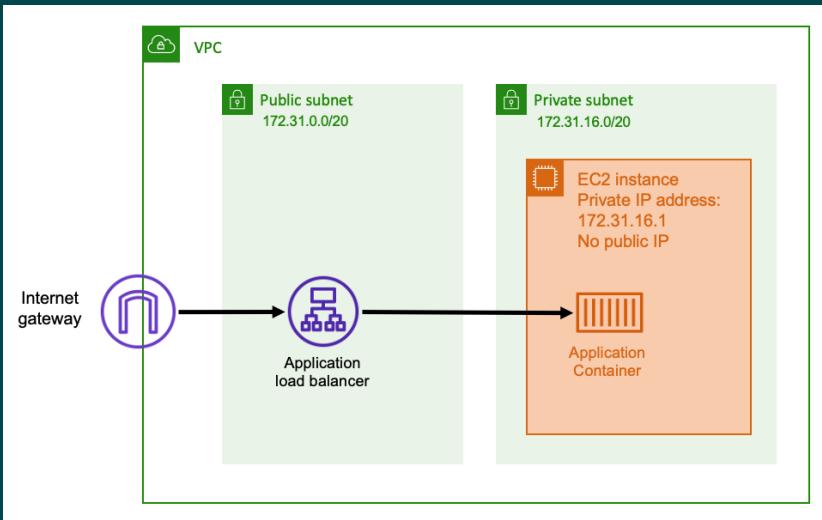


Private Subnet + NAT Gateway

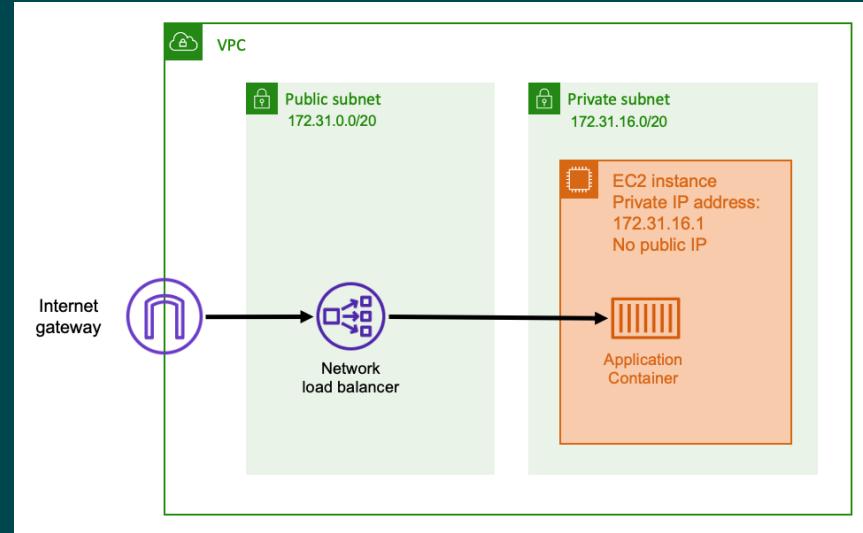
- Use case: when the application containers should be protected from external access
- E.g.: payment processing, sensitive data
- NAT gateway hourly charge applies



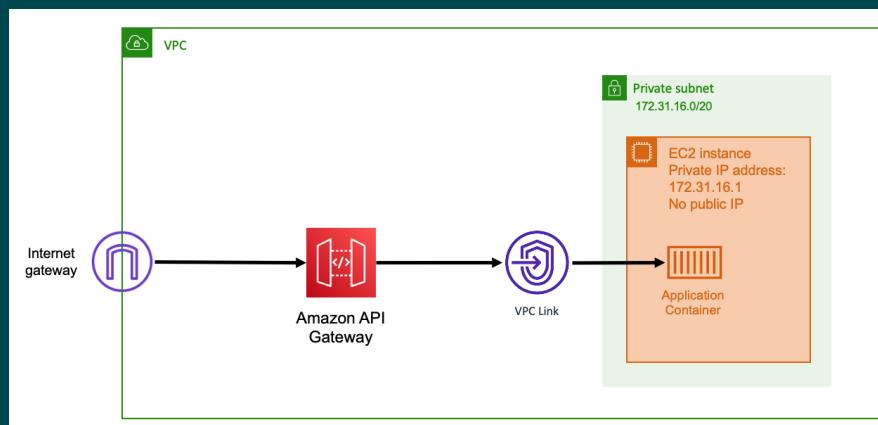
Receiving inbound connection from the internet



Application Load Balancer - for public HTTP services



Network Load Balancer – for non-HTTP protocols

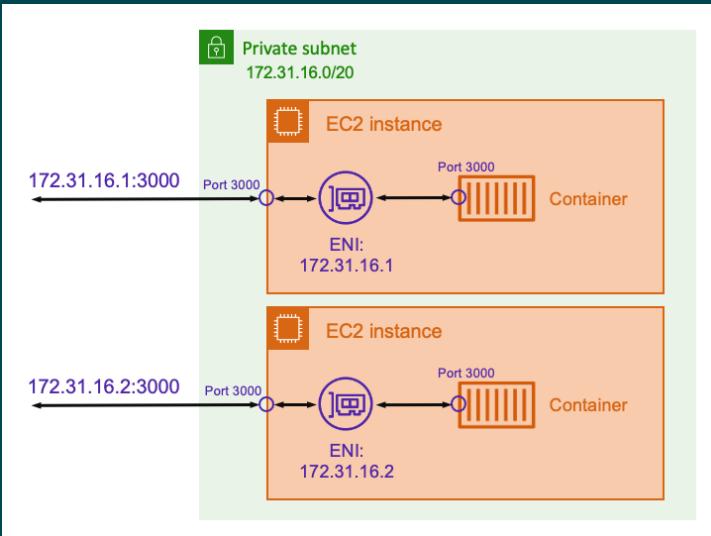


Amazon API Gateway - for HTTP applications with sudden bursts or low request volumes



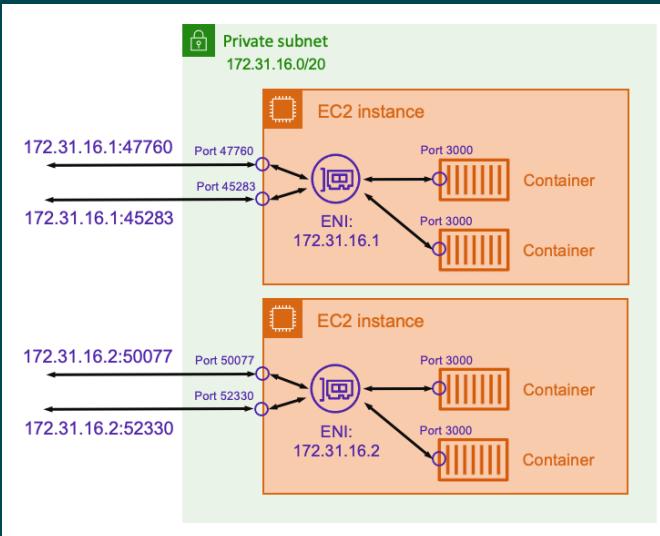
Choosing a network mode

Preventive Controls



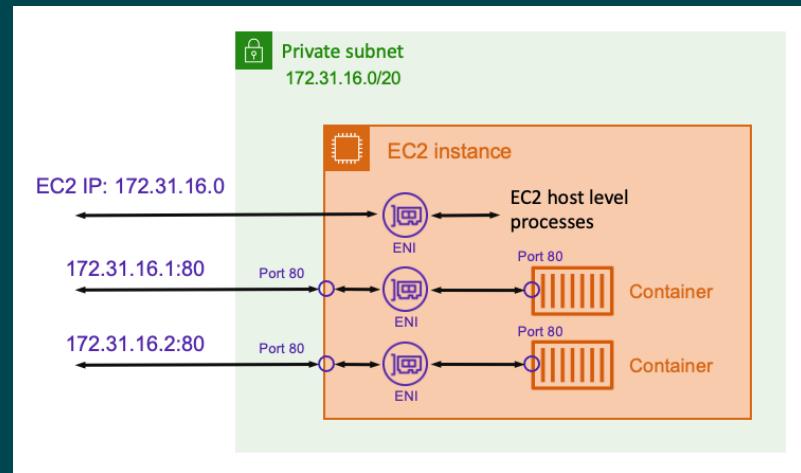
host mode

- container networking directly tied to the underlying host
- cannot run more than 1 task on an instance
- cannot remap port
- NOT RECOMMENDED



bridge mode

- a virtual network bridge between host and container networking
- supports dynamic port mapping
- difficult to lock down service to service communication
- not supported by ECS Fargate



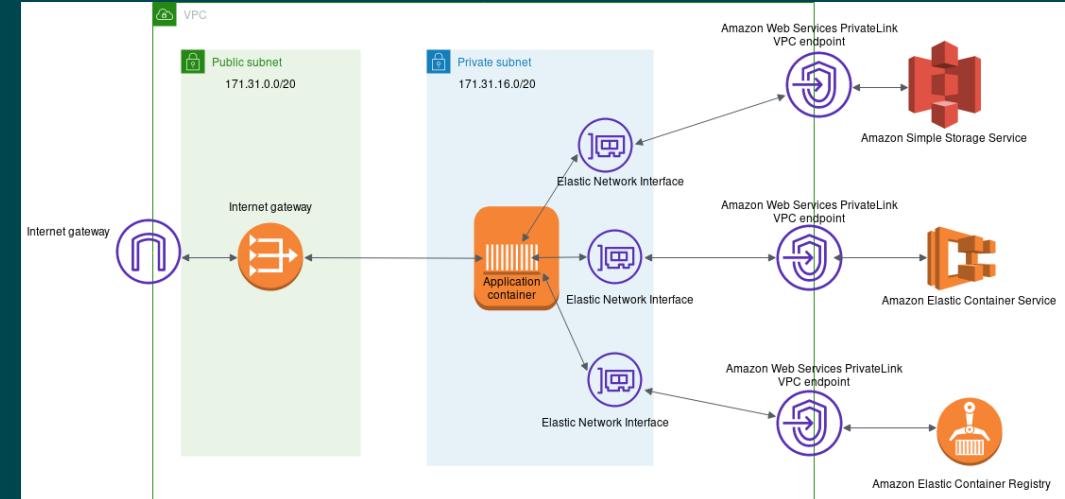
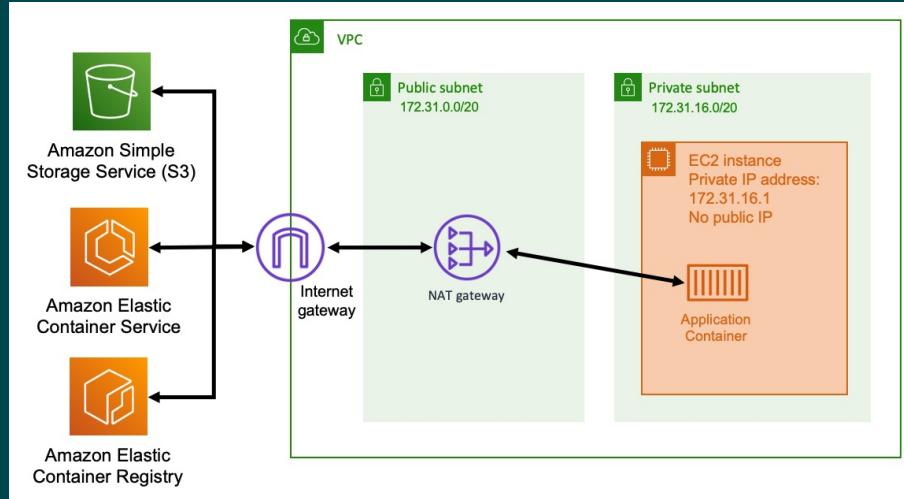
aws vpc mode

- each task gets its own IP
- tasks can bind to the same port
- better security group granularity
- IPv6 dual stack VPC compatible
- required when using Fargate
- considerations:
 - increase task density by ENI trunking
 - prevent IP address exhaustion



Connecting to AWS services from inside the VPC

Preventive Controls



NAT Gateway

- go through the public internet
- disadvantages
 - no control over outbound destinations
 - NAT gateway bandwidth limit

VPC Endpoint (AWS PrivateLink)

- Common
- private connectivity: traffic doesn't leave the AWS network

IAM best practice with ECS

Preventive Controls

- Block access to EC2 Metadata

```
sudo yum install -y iptables-services; sudo iptables --insert FORWARD 1 --in-interface docker+ --destination 192.0.2.0/32 --jump DROP
```

- Use “awsvpc” mode

- Use IAM access advisor to refine roles

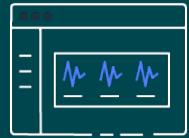
```
aws iam generate-service-last-accessed-details --arn arn:aws:iam::123456789012:policy/ExamplePolicy1
```

- Monitor AWS Cloudtrail for any suspicious activity



Amazon Inspector

Preventive Controls



Gain centralized visibility

- Environment coverage
- High impact findings
- Resources by finding severity



One-click continuous monitoring

- Automatic discovery of resources
- Monitors throughout the resource life-cycle



Prioritize with contextualized scoring

- Inspector Risk score
- Security metrics
- Customized views



Centrally manage at scale

- AWS Organizations
- Package vulnerability, Network reachability
- Environment coverage



Automate and take actions

- Management APIs
- Detailed findings in Amazon EventBridge
- Security Hub integration



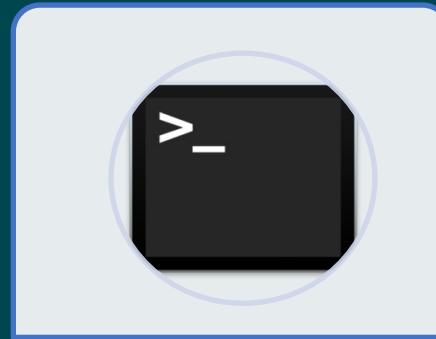
Prioritize vulnerability remediation

Types of vulnerability



Remotely
exploitable

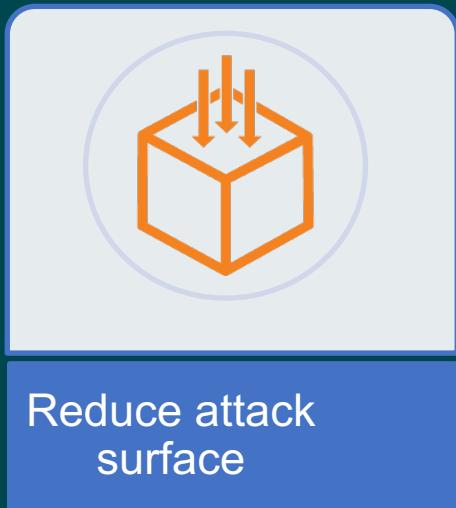
Remote Command/Code
Execution
SQL Injection
XSS
CSRF
Partial DDoS



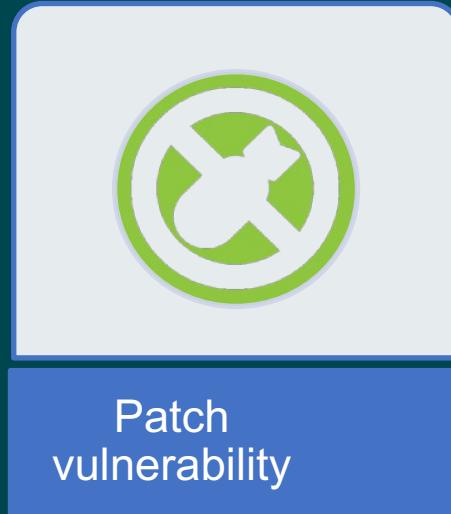
Locally
exploitable

Buffer overflow
Privilege escalation
Arbitrary execution
File and memory view as
root
System-wide DoS

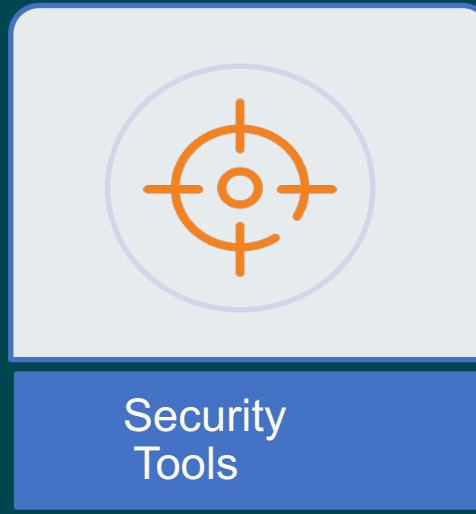
Vulnerability management



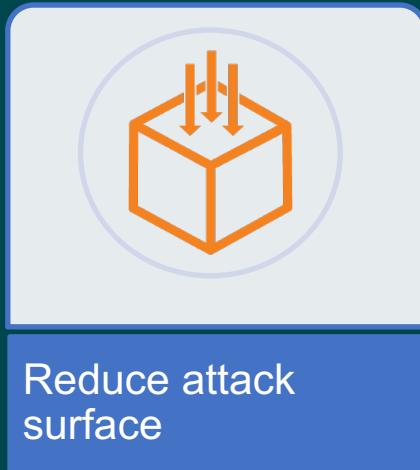
Reduce attack
surface



Patch
vulnerability

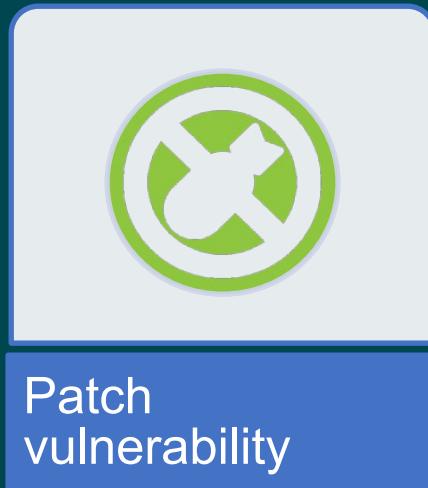


Security
Tools



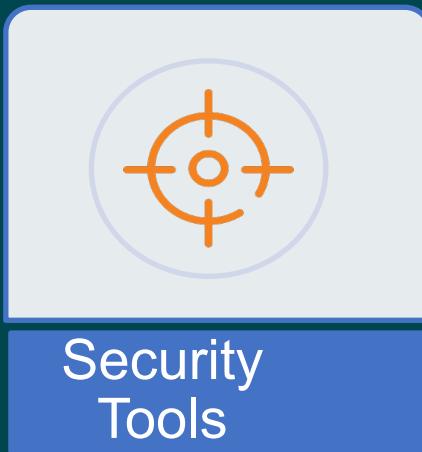
Reduce attack
surface

- Choose base image based wisely
 - Ubuntu vs Alpine (choose based on least vulnerability found)
 - Footprint and community response
 - Remove Bash (if not required)
 - Use apt-get command with --no-install-recommends
 - Build image from scratch with least amount of packages only required by OS
 - Use low level build tools like “buildah” for least attack surface



Patch
vulnerability

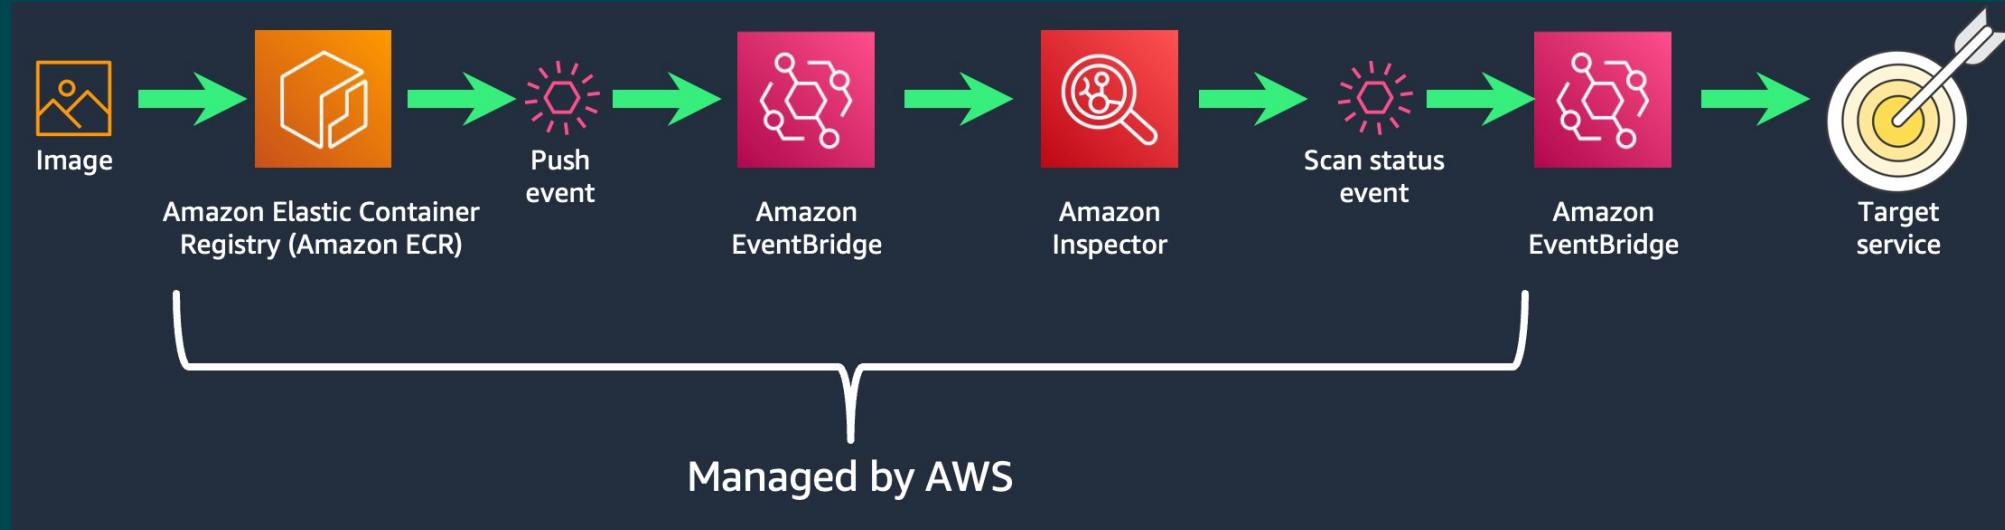
- Choose way to deploy application
 - Package manager
 - Use apt-get, yum, pip, npm etc. to update packages
 - Install from source
 - Most flexible and complex way to update dependent packages.



- Protect at Host, container and application
 - IPS (intrusion prevention system) for virtual patch
 - Malware detection
 - Support not just for Kubernetes but all containers
 - Admission control to reduce attack surface
 - RASP for protecting application inside out

Prevent vulnerable image to be deployed in production

Preventive Controls



- Use Amazon Inspector to find vulnerabilities in container images in elastic container registry and operating systems hosting containers on Amazon EC2
- Inspector leverages a continuous vulnerability scanning model to enable customers to support automated vulnerability analysis across all workloads

Other measures during build time

Preventive controls

Add Static code analysis in CI/CD pipeline

Add Container image scanning in CI/CD pipeline

- Find **Secrets** stored in Image
- Verify if all images have “**update**” statement.
- Secure DevOps **tool chain**
- **Vulnerabilities** risk at each layer in image
- Check for **malicious payload** in image
- Check for **compliance** in image (hardening best practice)
- Sign approved images at pipeline

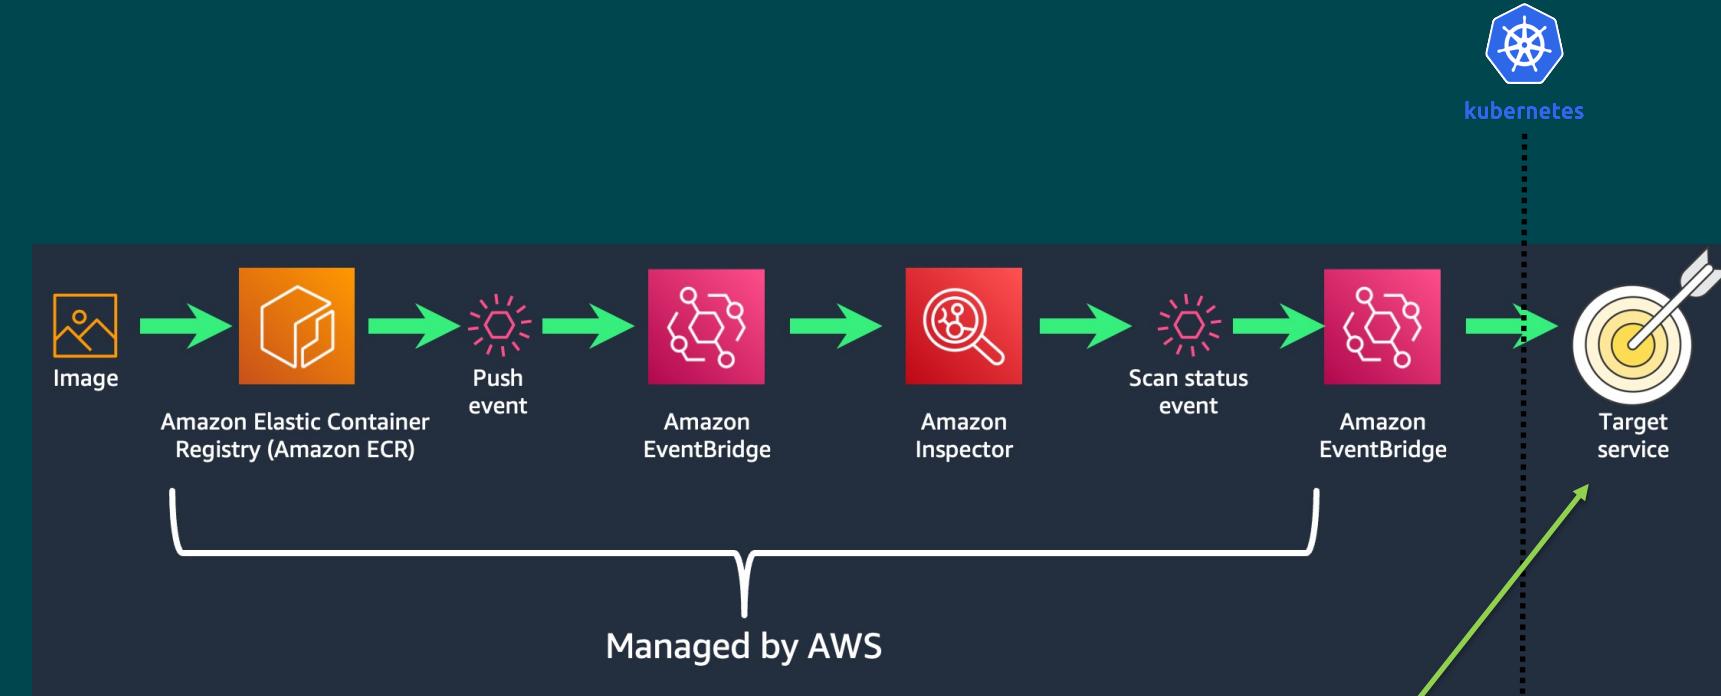


**Ensuring your Images are safe
for runtime**



Escape from DevOps pipeline

Preventive controls



```
dmathieu@: kubectl get pods --watch
NAME                                         READY   STATUS    RESTARTS   AGE
static-app-deployment-584d9757f5-4mxfq      1/1    Running   0          2m
static-app-deployment-584d9757f5-blwhr       1/1    Running   0          1m
static-app-deployment-584d9757f5-fm4f4       1/1    Running   0          1m
static-app-deployment-584d9757f5-fwnkf       1/1    Running   0          2m
static-app-deployment-584d9757f5-gbdtj       1/1    Running   0          1m
static-app-deployment-584d9757f5-n7fj4       1/1    Running   0          2m
static-app-deployment-584d9757f5-pnjgs       1/1    Running   0          1m
static-app-deployment-584d9757f5-r689q       1/1    Running   0          1m
static-app-deployment-584d9757f5-tck5         1/1    Running   0          56s
static-app-deployment-584d9757f5-zwt7x       1/1    Running   0          1m
static-app-deployment-5f6959d966-8lbqz        0/1    Pending   0          0s
static-app-deployment-5f6959d966-8lbqz        0/1    Pending   0          0s
static-app-deployment-5f6959d966-8lbqz        0/1    ContainerCreating 0          0s
```



Curse of privileged containers

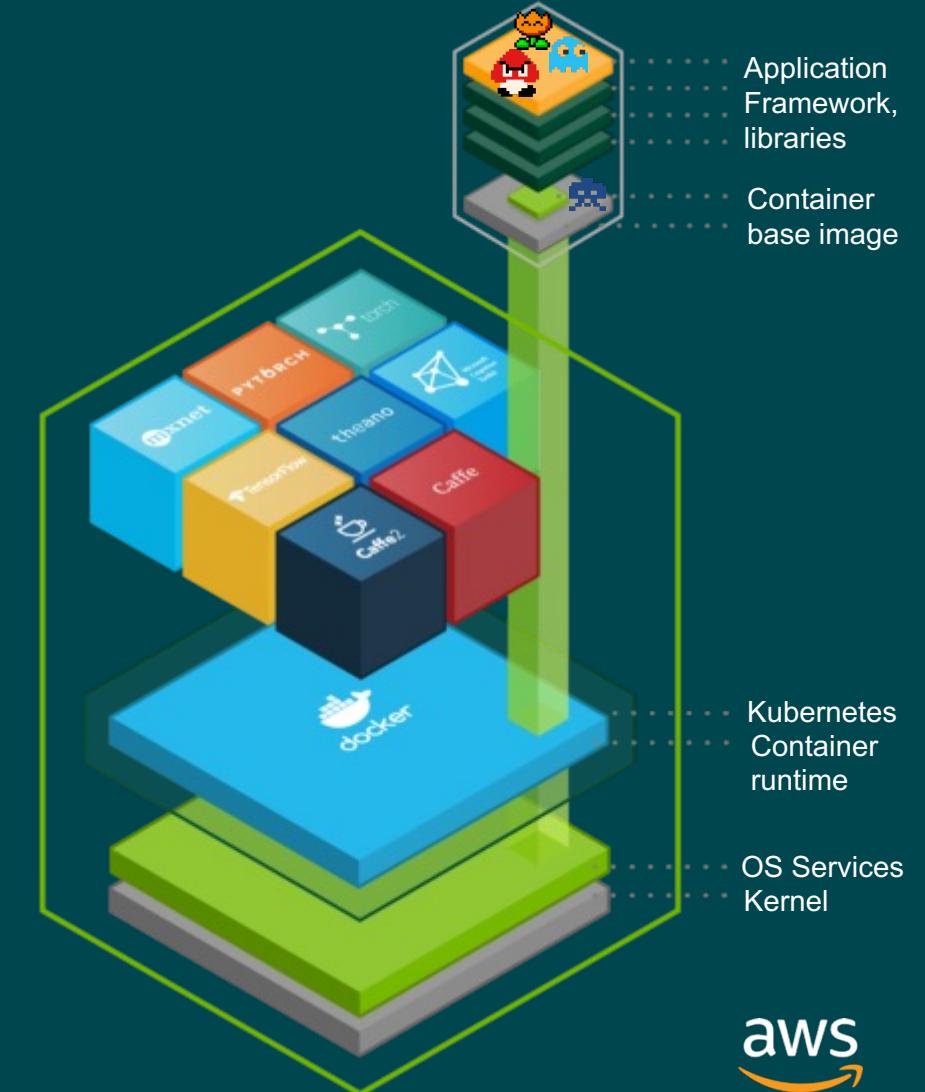
Preventive Controls

Container that is owned by Root

Vulnerable unpatched container platform

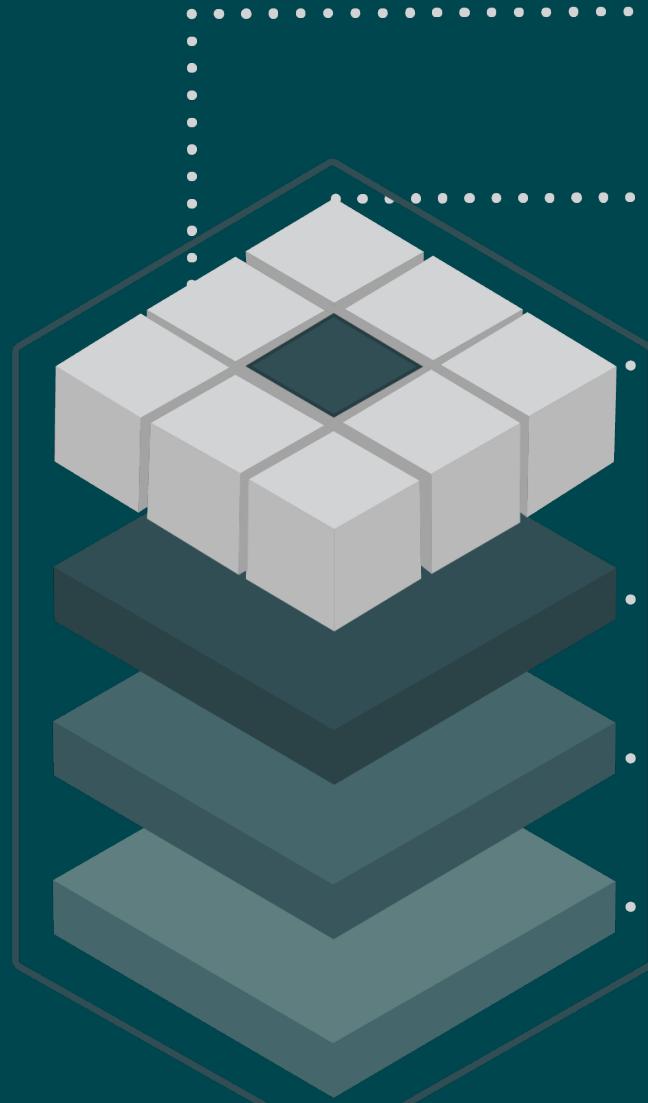
Container running in Host namespace

Containers with privilege escalation
enabled in K8S



Admission Control policy at runtime

Preventive Controls



..... Policy for prevent non-compliant runtime configuration

..... Policy for approved source of container registry

..... Container image signing for approved SBOM

..... Read-only container filesystem

..... Container node Security posture (e.g. CIS benchmark)

..... Select Container Optimized AMI



Do your customer ask you if their Container workload is secured and what are the ways to protect this containers ?

- Always
- Sometimes
- Never

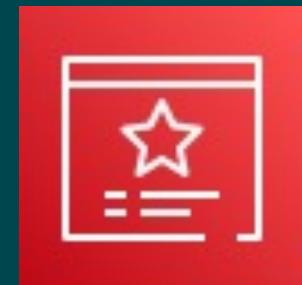


Encryption in Transit using ACM PCA

Preventive Controls

ACM Private CA Kubernetes Connector

- Easily automate and configure end-to-end encryption for Amazon Elastic Kubernetes Service (EKS)
- Generate and manage private certificates for TLS and mTLS for Kubernetes clusters
- Simplify regulatory and compliance requirements, such as PCI-DSS and HIPAA, for data in-transit



Ingress Controllers and Load Balancers

Ingress controllers are a way for you **to intelligently route HTTP/S traffic** that emanates from outside the cluster to services running inside the cluster.

Ingresses are often fronted by a Network Load Balancer (**NLB**) or an Application Load Balancer (**ALB**)

Encrypted traffic can be terminated at different places within the network:

- **The pod** – Supports end-to-end encryption, but places additional burden on your Pod to establish TLS handshakes
- **The load balancer** – Reduces Pod compute cycles needed to establish TLS handshake
- **Ingress resource** – Reduces Pod compute cycles needed to establish TLS handshake



Elastic Load Balancing



Application Load
Balancer



Network Load
Balancer



Encrypting Kubernetes Secrets

Preventive Controls

There are two primary options for protecting secrets in EKS:

1. Use AWS KMS for envelope encryption of Kubernetes secrets
2. [Recommended] Use an external secrets management service – AWS Secrets Manager or a partner solution from the AWS Marketplace



AWS KMS



AWS Secrets Manager

Secrets on EKS: Envelope Encryption

Envelope Encryption on KMS – Key features

- The **KMS plugin for Kubernetes** allows you to encrypt your secrets with a unique **data key** generated by AWS KMS
- The **data key** is then encrypted using a **KMS key** which is non-exportable and cannot leave the hardware security modules that back KMS
- **KMS keys** are rotated on a recurring schedule (if rotation is enabled)
- Using the **KMS plugin for Kubernetes**, all secrets are stored (encrypted) in etcd and can only be decrypted by the Kubernetes API server



Secrets on EKS: AWS Secret Manager

Allows for automated secrets rotation

Kubernetes doesn't automatically rotate secrets. If you have to rotate secrets, consider using an external secret store, e.g. AWS Secrets Manager.

Improved security over default Kubernetes secrets management

AWS Secrets Manager offers features such as fine-grained access controls, encryption by default (cannot be disabled), and automatic rotation of secrets. None of these capabilities are available with default Kubernetes secrets management.

Secret Store CSI Driver for direct integration with AWS services

The [Secret Store CSI Driver](#) is a community project that uses the CSI driver model to fetch secrets from external secret stores.

The AWS provider for the Secret Store CSI Driver supports:

- [AWS Secrets Manager & AWS Systems Manager Parameter Store](#)
- Rotation of expired secrets
- Synchronization of AWS Secrets Manager secrets to Kubernetes Secrets - synchronization can be useful when you need to reference a secret as an environment variable instead of reading them from a volume



Additional best practices for secrets

Use separate namespaces as a way to isolate secrets from different applications

If you have secrets that cannot be shared between applications, create a separate namespace for those applications.

For example: corp/secrets/accounting vs. corp/secrets/engineering

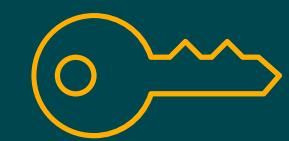
Use volume mounts instead of environment variables

The values of environment variables can unintentionally appear in logs. Secrets mounted as volumes are automatically removed from the node when the pod is deleted.

Audit the use of Kubernetes Secrets

Turn on EKS audit logging and create a CloudWatch metrics filter and alarm to alert you when a secret is used



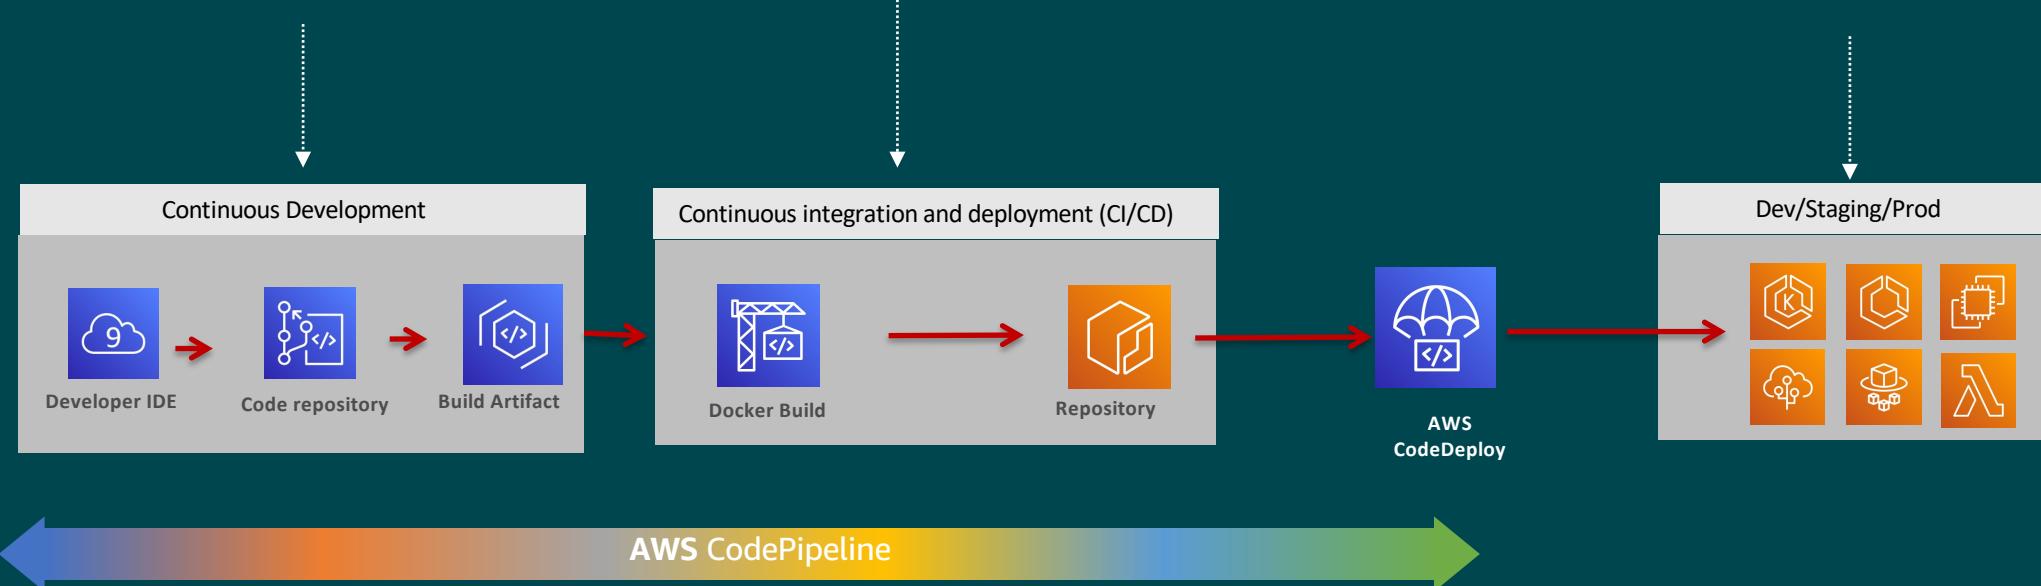


Preventative

Amazon CodeGuru
AWS Cloud9 IDE
AWS CodeArtifact
AWS CodeCommit

Amazon Inspector
AWS CodeBuild
AWS ECR

- AWS Certificate manager
- AWS Key Management service
- AWS Secrets manager
- AWS IAM
- AWS Systems manager
- AWS PrivateLink





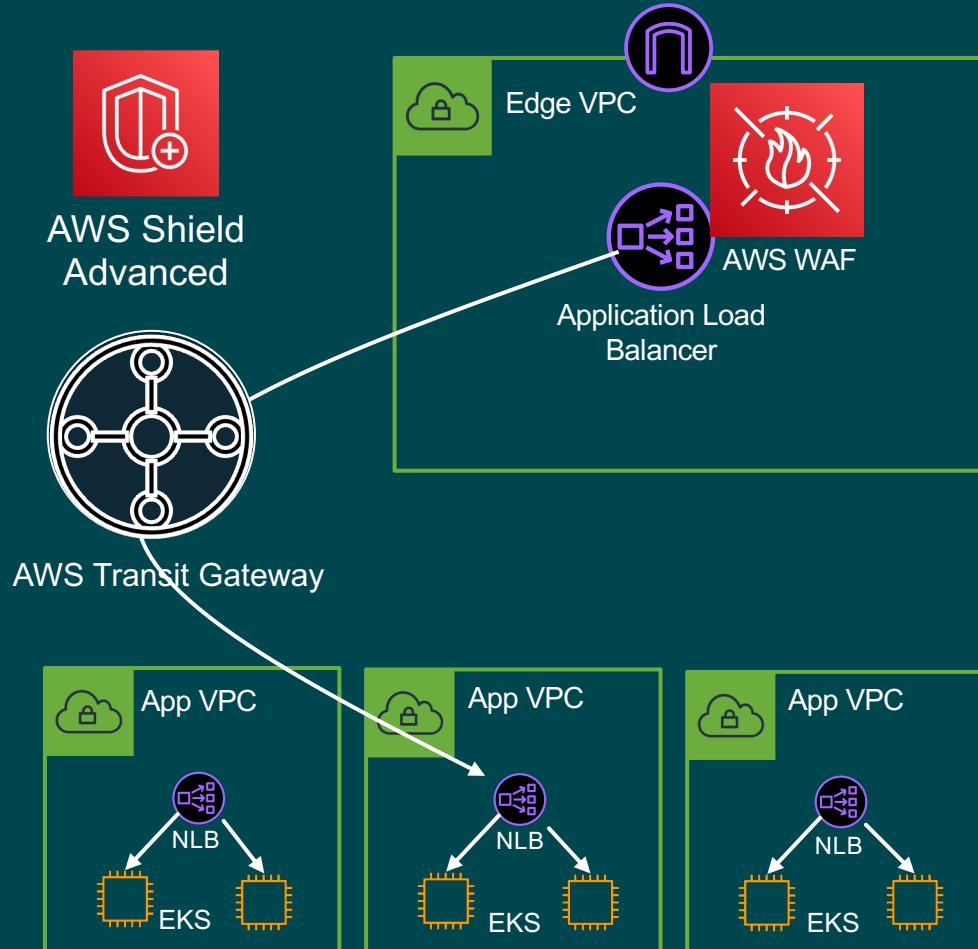
Demo



Protective controls

Protect container-based applications from attacks

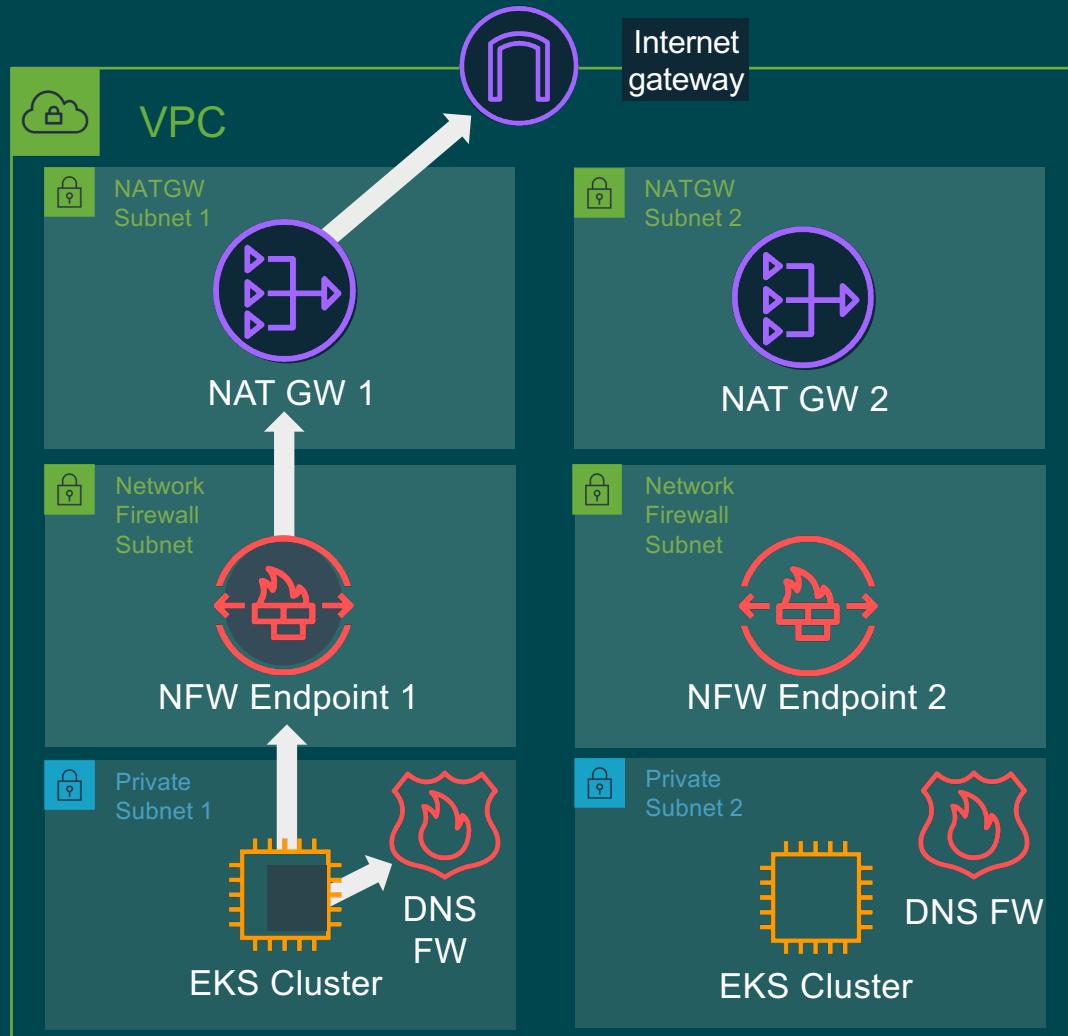
Edge Protection with AWS Shield Advanced and AWS WAF



- Stop DDoS attacks with AWS Shield Advanced and Automatic Application Layer(L7) Mitigation
- Protect against OWASP top 10 web attack such as SQL injection and Cross-Site Scripting
- Gain additional visibility and control over malicious bot traffic
- Support for centralized (shown) and distributed application architectures

Mitigate software supply-chain risk

Amazon Route 53 Resolver DNS Firewall and AWS Network Firewall

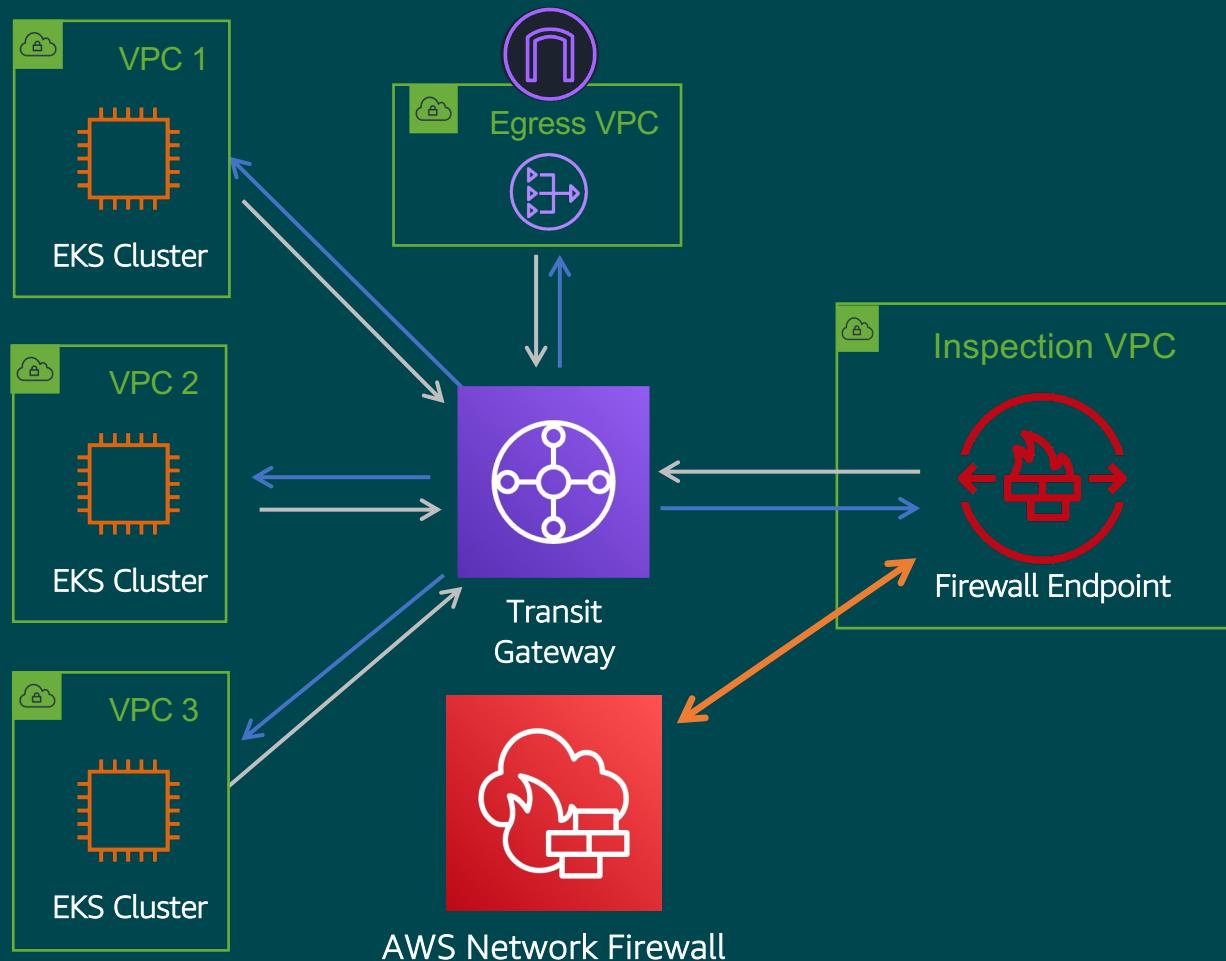


- Detect and block unexpected network egress communications from malicious code
- Limit network communications to trusted code and container repositories
- Add network protocol detection and enforcement to egress controls
- Support for centralized or distributed (shown) deployment models



Block malicious network activity

AWS Network Firewall Managed Threat Signatures



- AWS Network Firewall includes 16 threat signature categories including coin mining, web attacks and emerging events
- Full visibility into signature content to selectively choose what traffic to block or alert against
- Inspect traffic in centralized (shown) or distributed deployment models





Demo



Threat Detection and Response





THREAT Detection

Continuously check your AWS accounts and workloads to detect threats and verify compliance based on industry standards.



AWS Detective

Makes it easy to analyze, investigate, and quickly identify the root cause of potential security issues or suspicious activities



AWS GuardDuty

Threat detection service that continuously monitors for malicious activity and unauthorized behavior to protect your AWS accounts, workloads, and data stored in Amazon S3



AWS Inspector

Automated security assessment service that helps improve the security and compliance of applications deployed on AWS

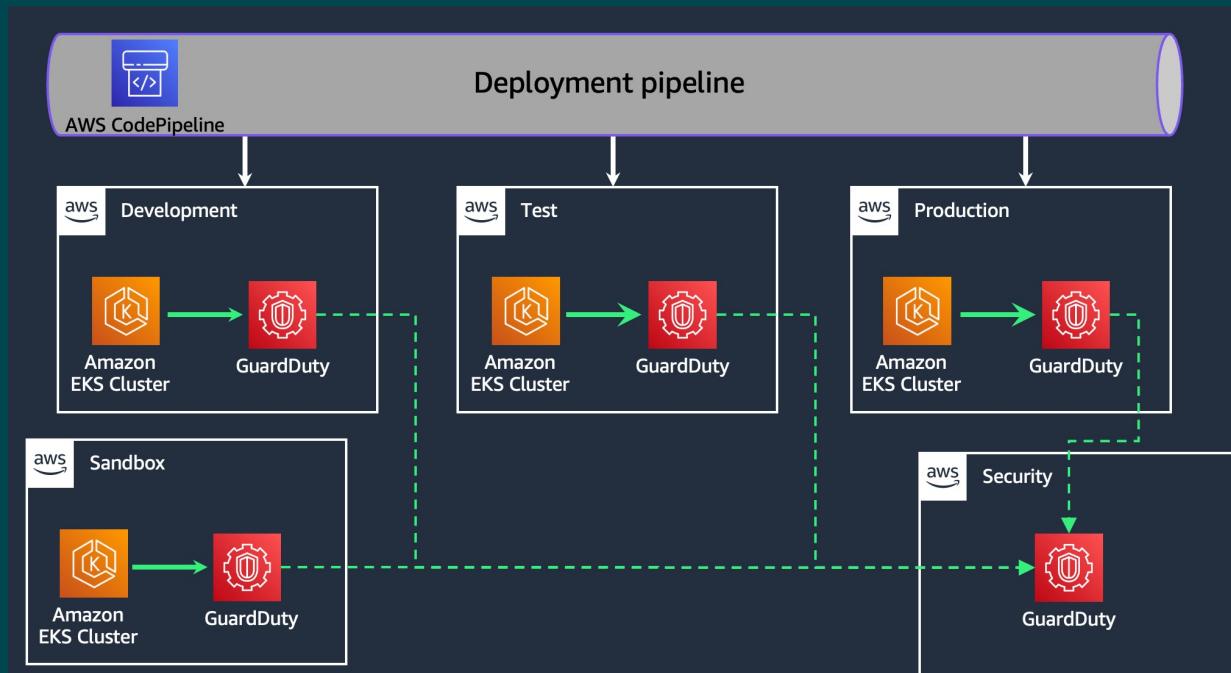


Amazon GuardDuty



Detect malicious activity running in container workloads

Threat Detection with Amazon GuardDuty



- Review Amazon EKS specific activity, such as pod volume patterns and container service user activity, including divergent behavior within and across EKS clusters
- Investigate security findings with EKS clusters, such as cryptocurrency mining, unintentional admin privilege exposure, container misconfigurations that allow access to underlying EC2 nodes, or behavioral patterns common to compromised container clusters

GuardDuty for Kubernetes Finding Types

Policy

- Exposed dashboard
- Admin access to default service account
- Anonymous access granted

Malicious access

- Data discovery, exfiltration, or modification from:
 - Tor
 - Successful anonymous access
 - Malicious IPs

Suspicious behavior

- Execution in Kubernetes system pod
- Container with sensitive mount
- Privilege container

- GuardDuty immediately begins to analyze Kubernetes data sources from your Amazon EKS clusters and monitors them for malicious and suspicious activity.

GuardDuty aligns findings using the MITRE ATT&CK framework

Credential Access Defense Evasion Discovery Impact
Privilege Escalation Policy Execution Persistence



Amazon GuardDuty Malware Protection

- **GuardDuty Malware Protection** – a fully managed malware detection service supports detection of malicious software by scanning Amazon Elastic Block Storage (EBS) and container workloads running on Amazon EC2.

- Detects malware (trojans, worms, rootkits, crypto miners, bots, etc.) and generates a new finding in GuardDuty
- Scans **Elastic Block Storage (EBS)** for malware on Amazon EC2 instance and container workloads that are exhibiting suspicious behavior
 - *Automatic - malware scanning triggered on GuardDuty findings*
 - *Agentless - no security software required to install or maintain*



GuardDuty Malware Protection

DELIVERS AGENTLESS DETECTION OF MALWARE ON AWS WORKLOADS



Single-click
organization-wide
malicious file detection



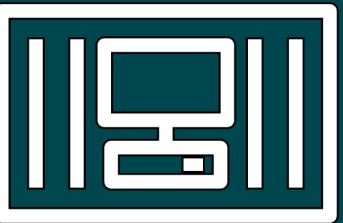
Centralized monitoring,
automation, and investigation



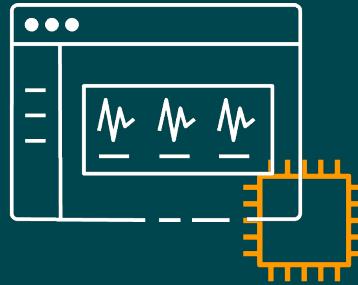
Contextualized
findings to validate
suspicious behavior



No agents to install,
update, or maintain



Container aware



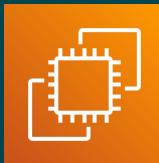
No performance impact or
hidden Amazon EC2
costs for scanning



GuardDuty Malware Protection: Coverage

Amazon EC2 coverage

Amazon EC2 instances



Amazon ECS



Amazon EKS



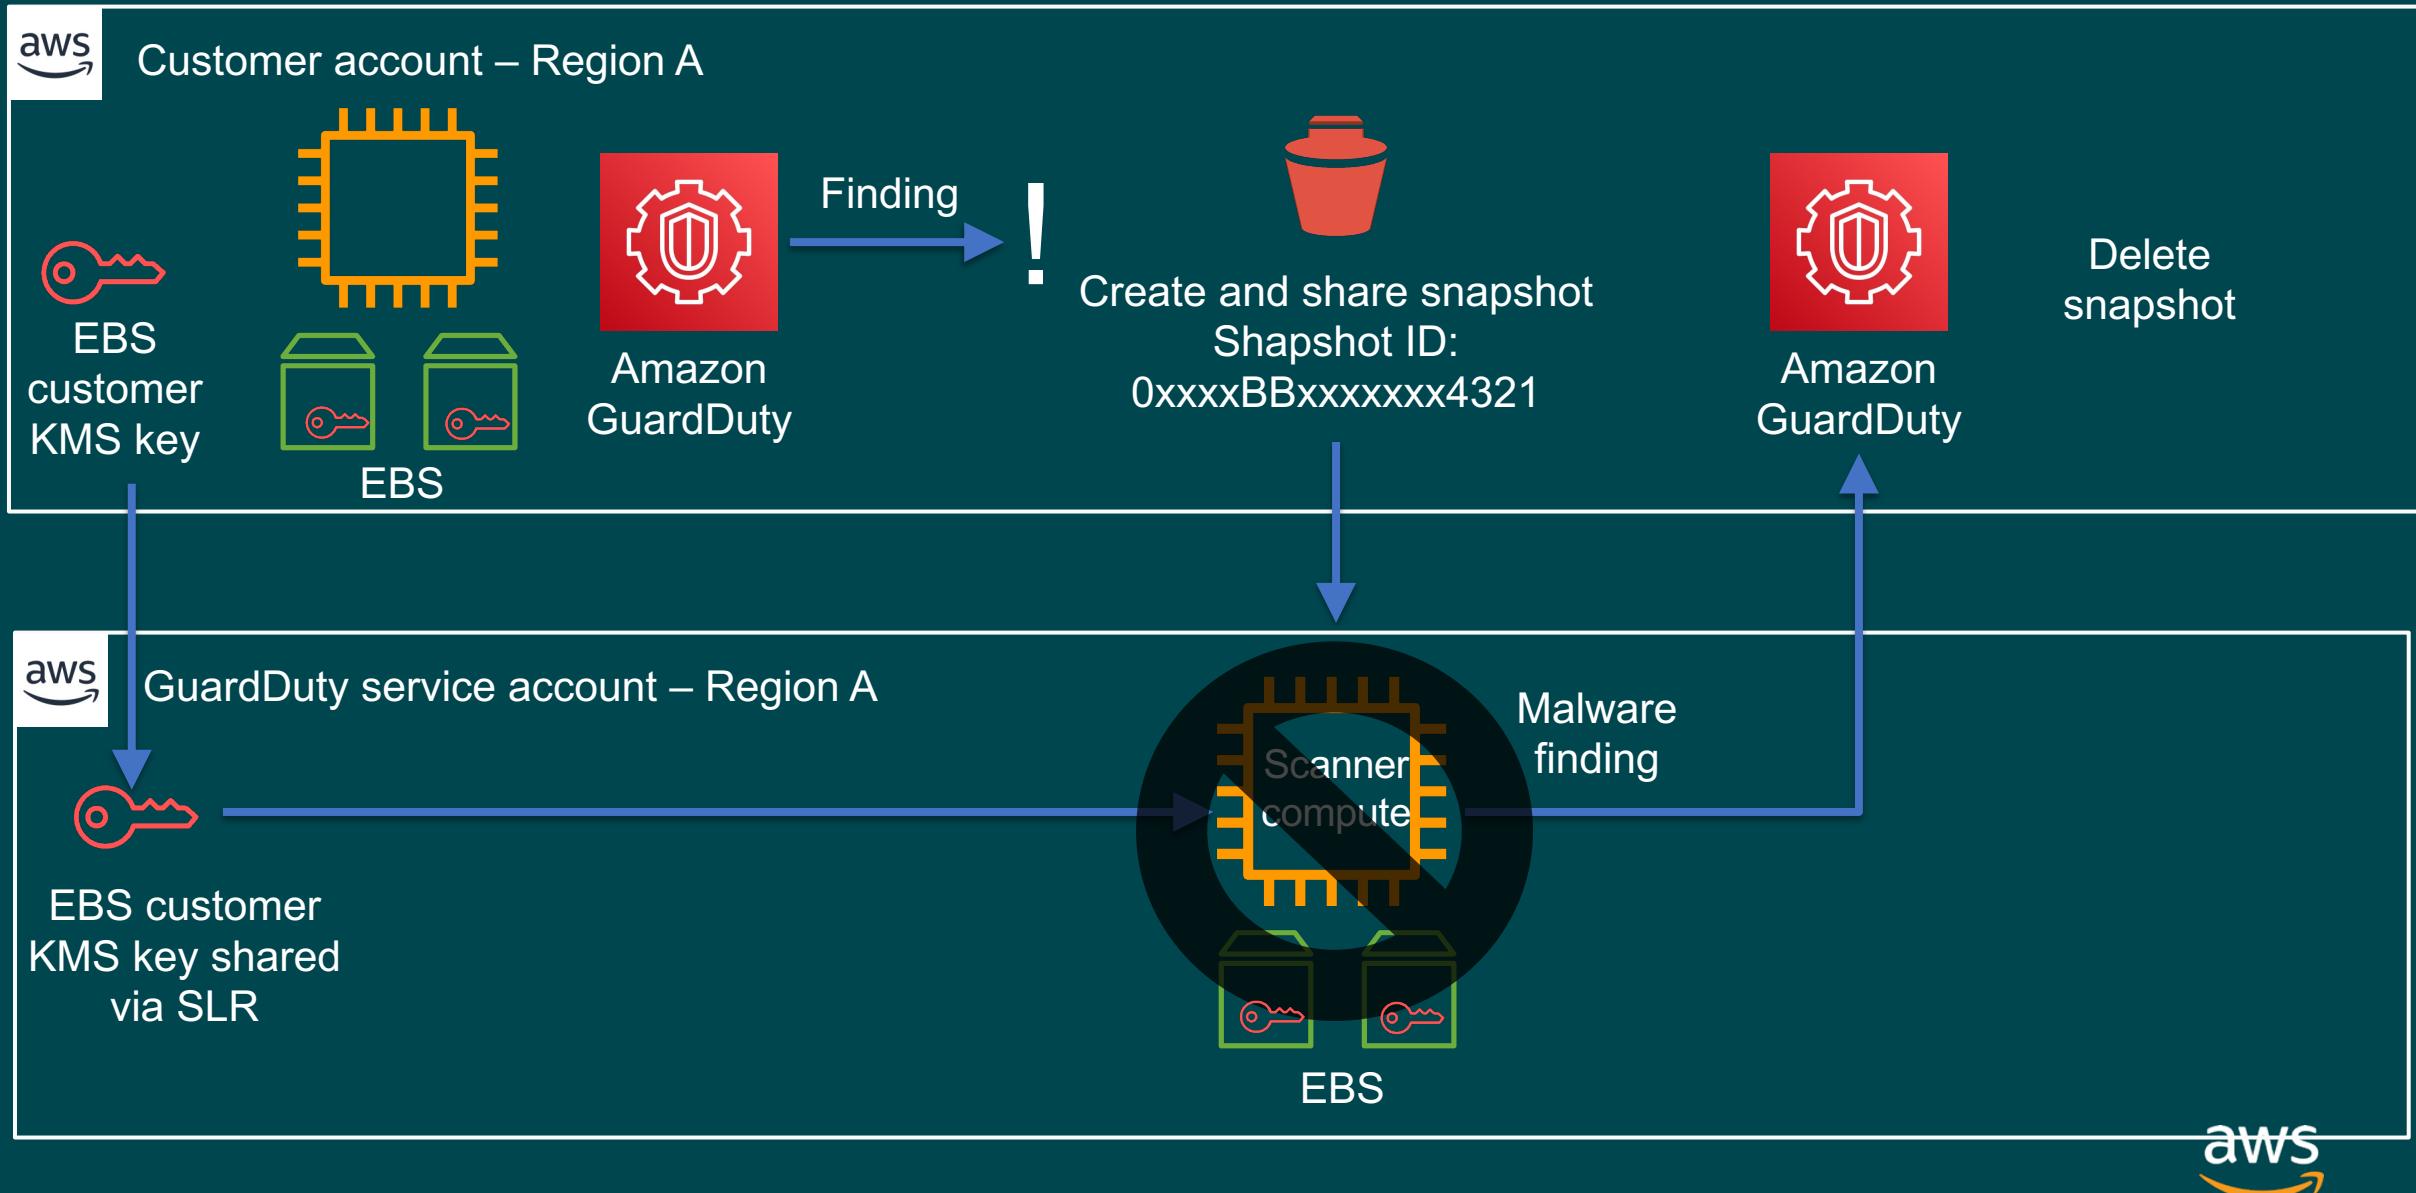
Self-managed containers
on Amazon EC2



When does GuardDuty initiate a malware scan?

- Malware scans are automatically triggered when GuardDuty detects a potentially compromised Amazon EC2 instance to identify malware that may be causing the activity
- It only scans an EC2 instance once every 24 hours, irrespective of multiple GuardDuty findings observed on it

How does it work?



Investigate events



Investigate security issues in container workloads

Identify the root-cause of malicious or suspicious behavior with Amazon Detective



- Review Amazon EKS specific activity, such as pod volume patterns and container service user activity, including divergent behavior within and across EKS clusters
- Investigate security findings with EKS clusters, such as cryptocurrency mining, unintentional admin privilege exposure, container misconfigurations that allow access to underlying EC2 nodes, or behavioral patterns common to compromised container clusters

Amazon Detective Use Case Benefits

INVESTIGATE ISSUES FASTER AND WITH LESS EFFORT



1) Finding / Alert triage

Accelerate triage and avoid unnecessary escalations



2) Incident Investigation

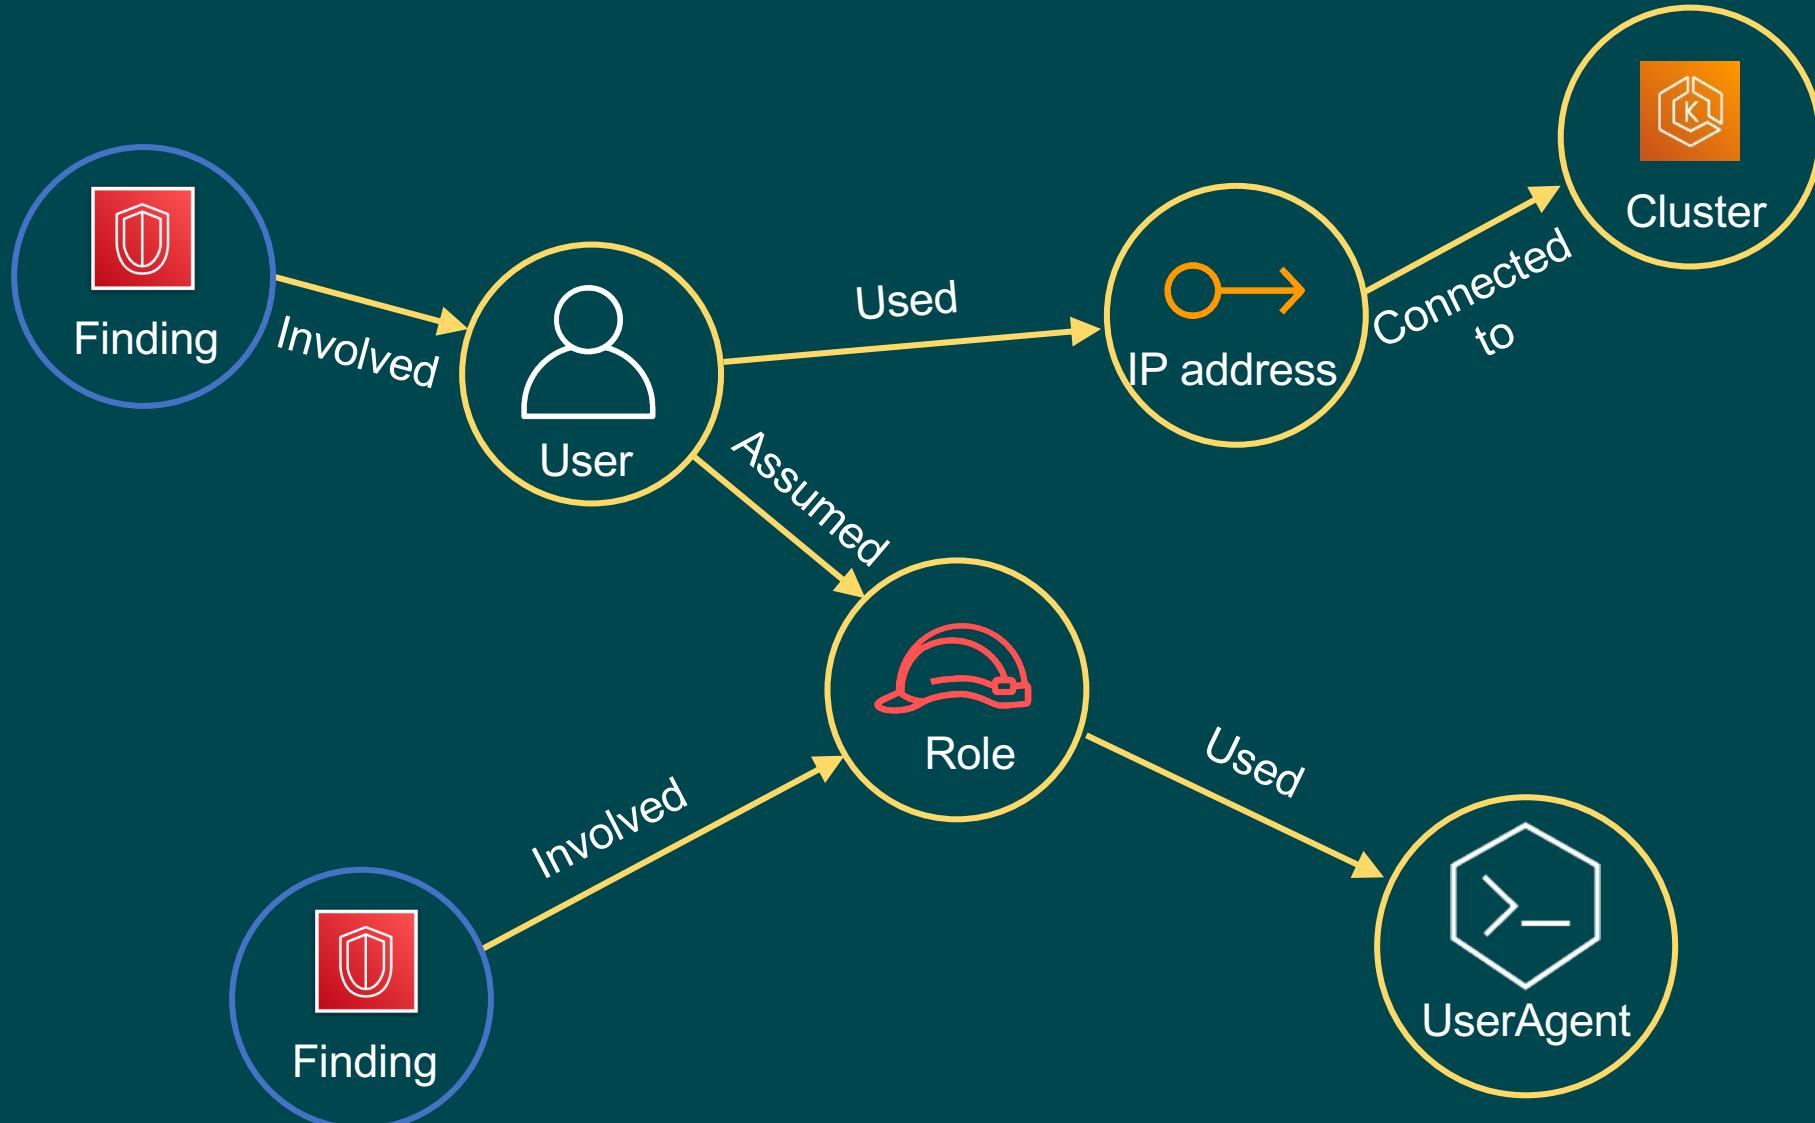
Improve context and surface correlated behavior



3) Threat Hunting

Simplify data collection, aggregation and pivoting

Amazon Detective - Security behavior graph



Amazon Detective support for EKS

- **Amazon Detective** security investigations for Amazon Elastic Kubernetes Service (Amazon EKS) clusters to quickly **analyze**, **investigate**, and identify the **root-cause** of malicious or suspicious behavior that represents potential threats to container workloads.
- Review **Amazon EKS** specific activity, such as **pod volume patterns** and **container service user activity**, including divergent behavior within and across EKS clusters
- Investigate security findings with their **EKS clusters**, such as cryptocurrency mining, unintentional **admin privilege exposure**, container **misconfigurations** that allow access to underlying EC2 nodes, or behavioral patterns common to **compromised** container clusters.



Demo



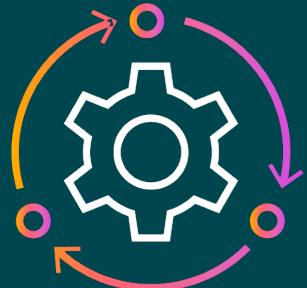
Do you have customers asking for container runtime protection?

- Yes
- No

Continuous monitoring with automated response & remediation



Security Hub – response & remediation



Enable **automated remediation** for high-severity configuration findings



Use **custom actions** to invoke runbooks for automated response



Use **partner integrations** to consolidate and normalize security findings



Integrate with **ticketing and workflow** tools



Simple and Scalable Security Monitoring

Scale existing services

customers use same console, findings, and experience



Amazon GuardDuty



Simple and easy deployment
One-click enables container support
AWS Orgs assures environment-wide enablement
(new customers have support on by default)



Amazon Detective



Detective EKS Investigation
Detective is container-aware and continuously aggregates KAL into graph model and analytics. Pivot from GuardDuty console to immediately investigate findings for root cause analysis.



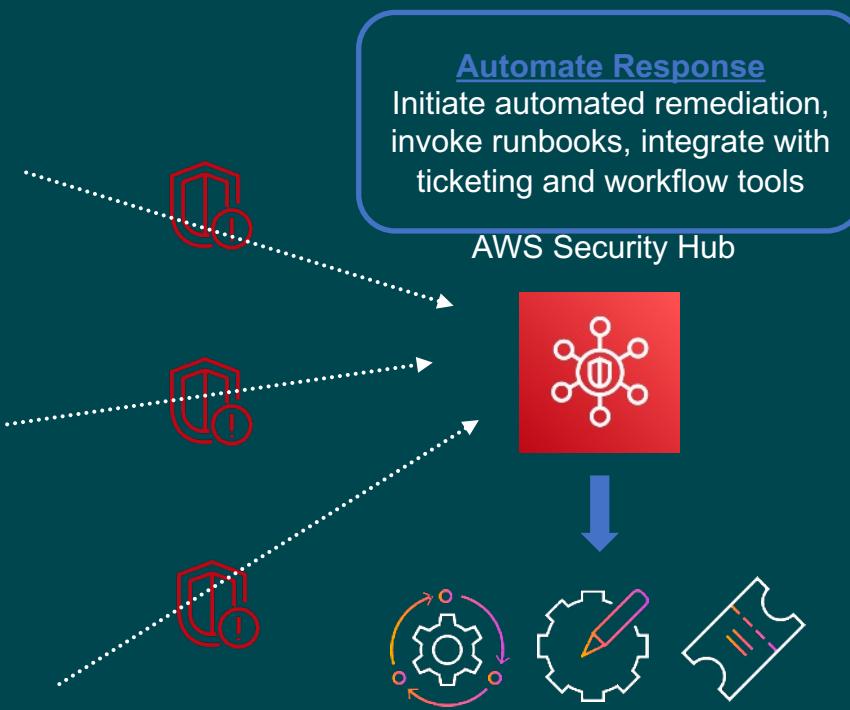
Amazon Inspector



Inspector ECR Support
Inspector scans ECR images on push and continually for software vulnerability management—integrates with ECR console for easy builder communication

Continuous monitoring

Centralization of security findings scales and automates operations



Automate and reduce risk with integrated services

Comprehensive set of APIs
and security tools



Continuous monitoring
and protection



Threat remediation
and response



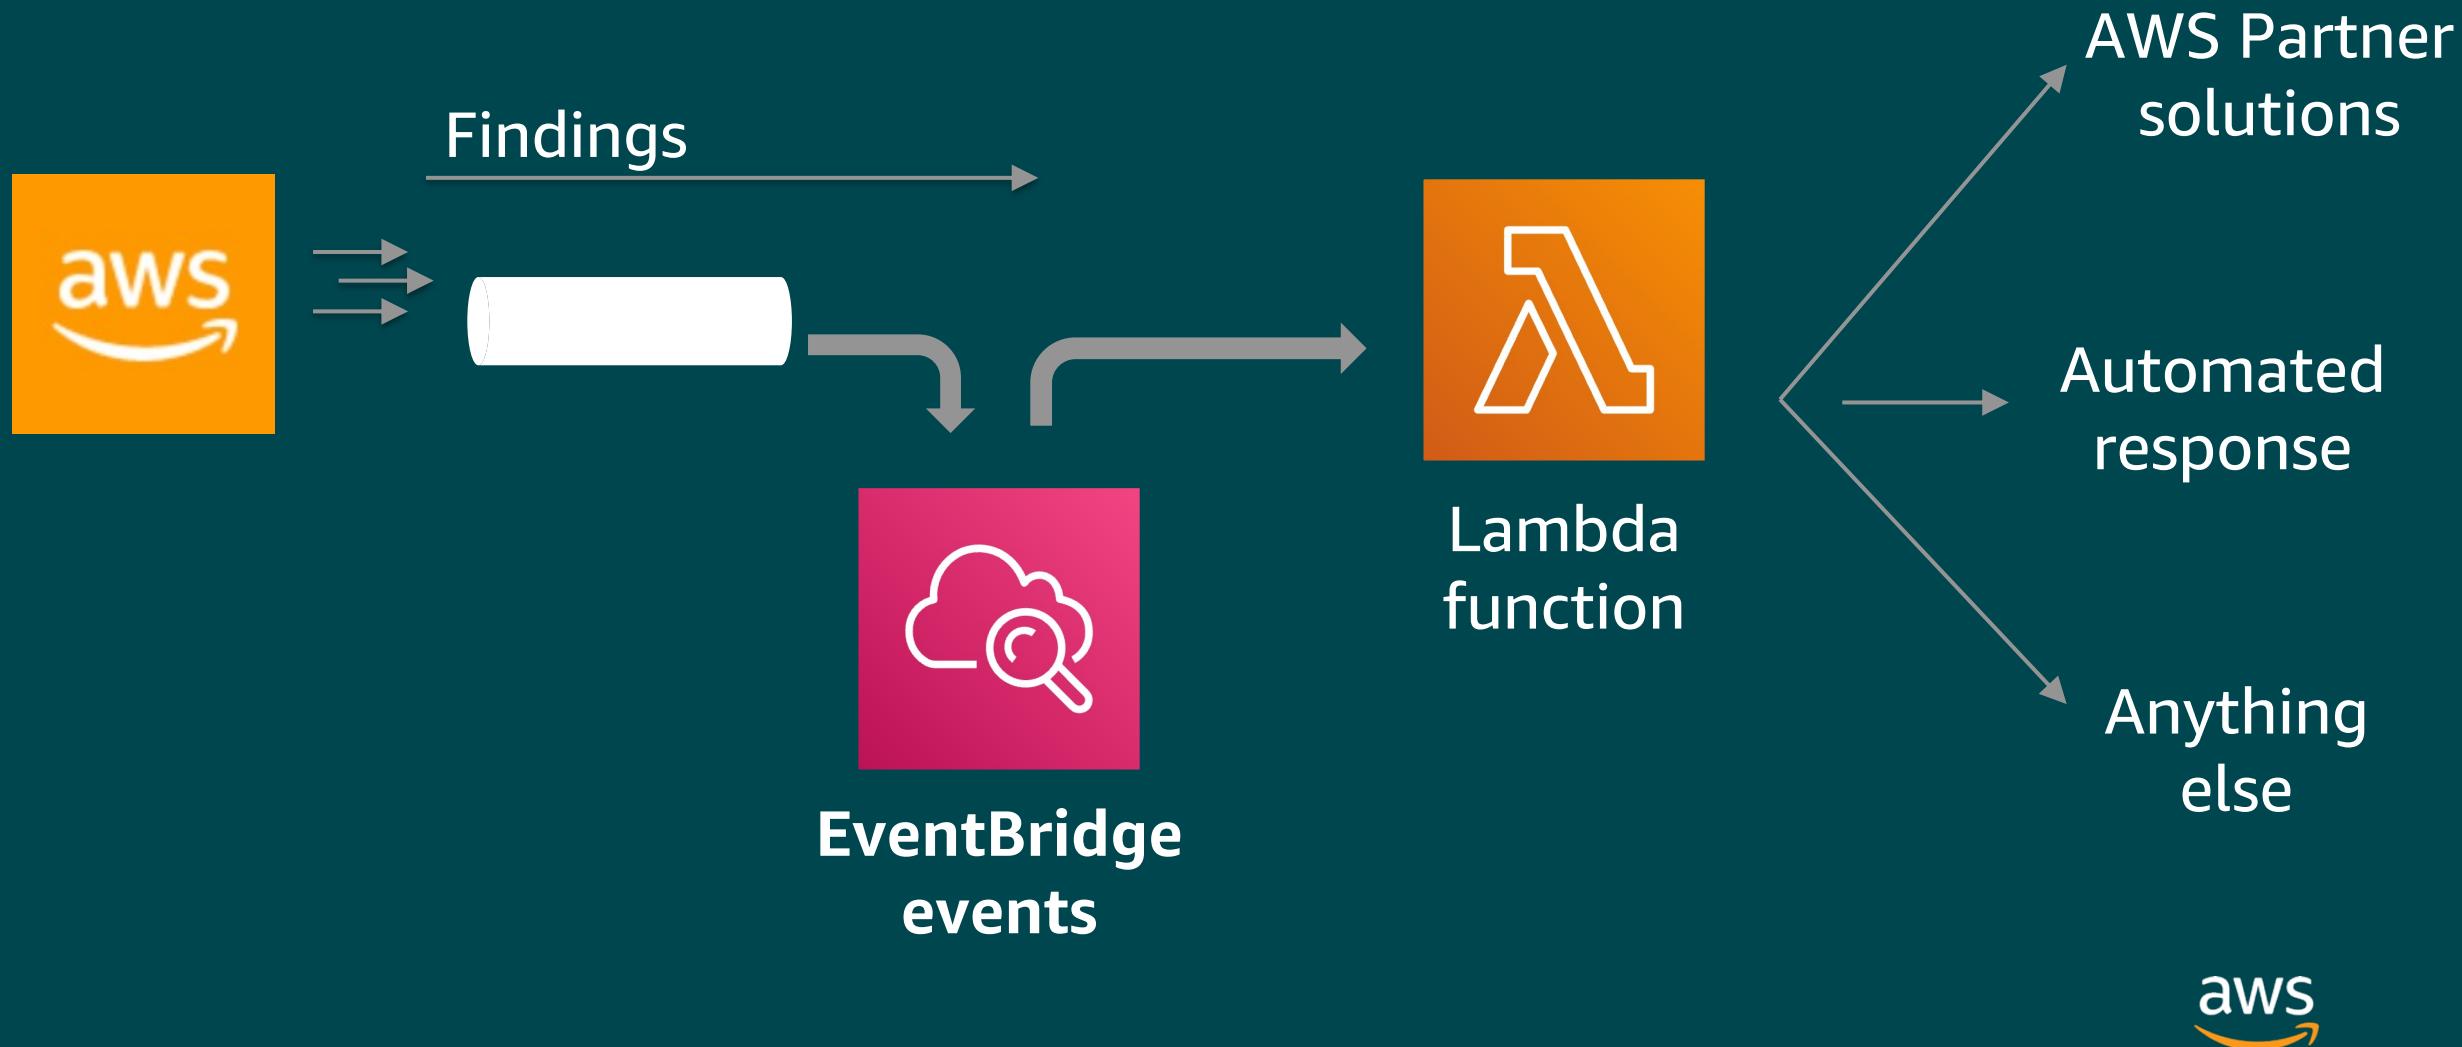
Operational efficiencies to
focus on critical issues



Securely deploy business
critical applications



Threat response: High-level playbook



Threat response: Services



Lambda

Run code for virtually
any kind of
application or
backend service –
zero administration



Systems Manager



Amazon VPC

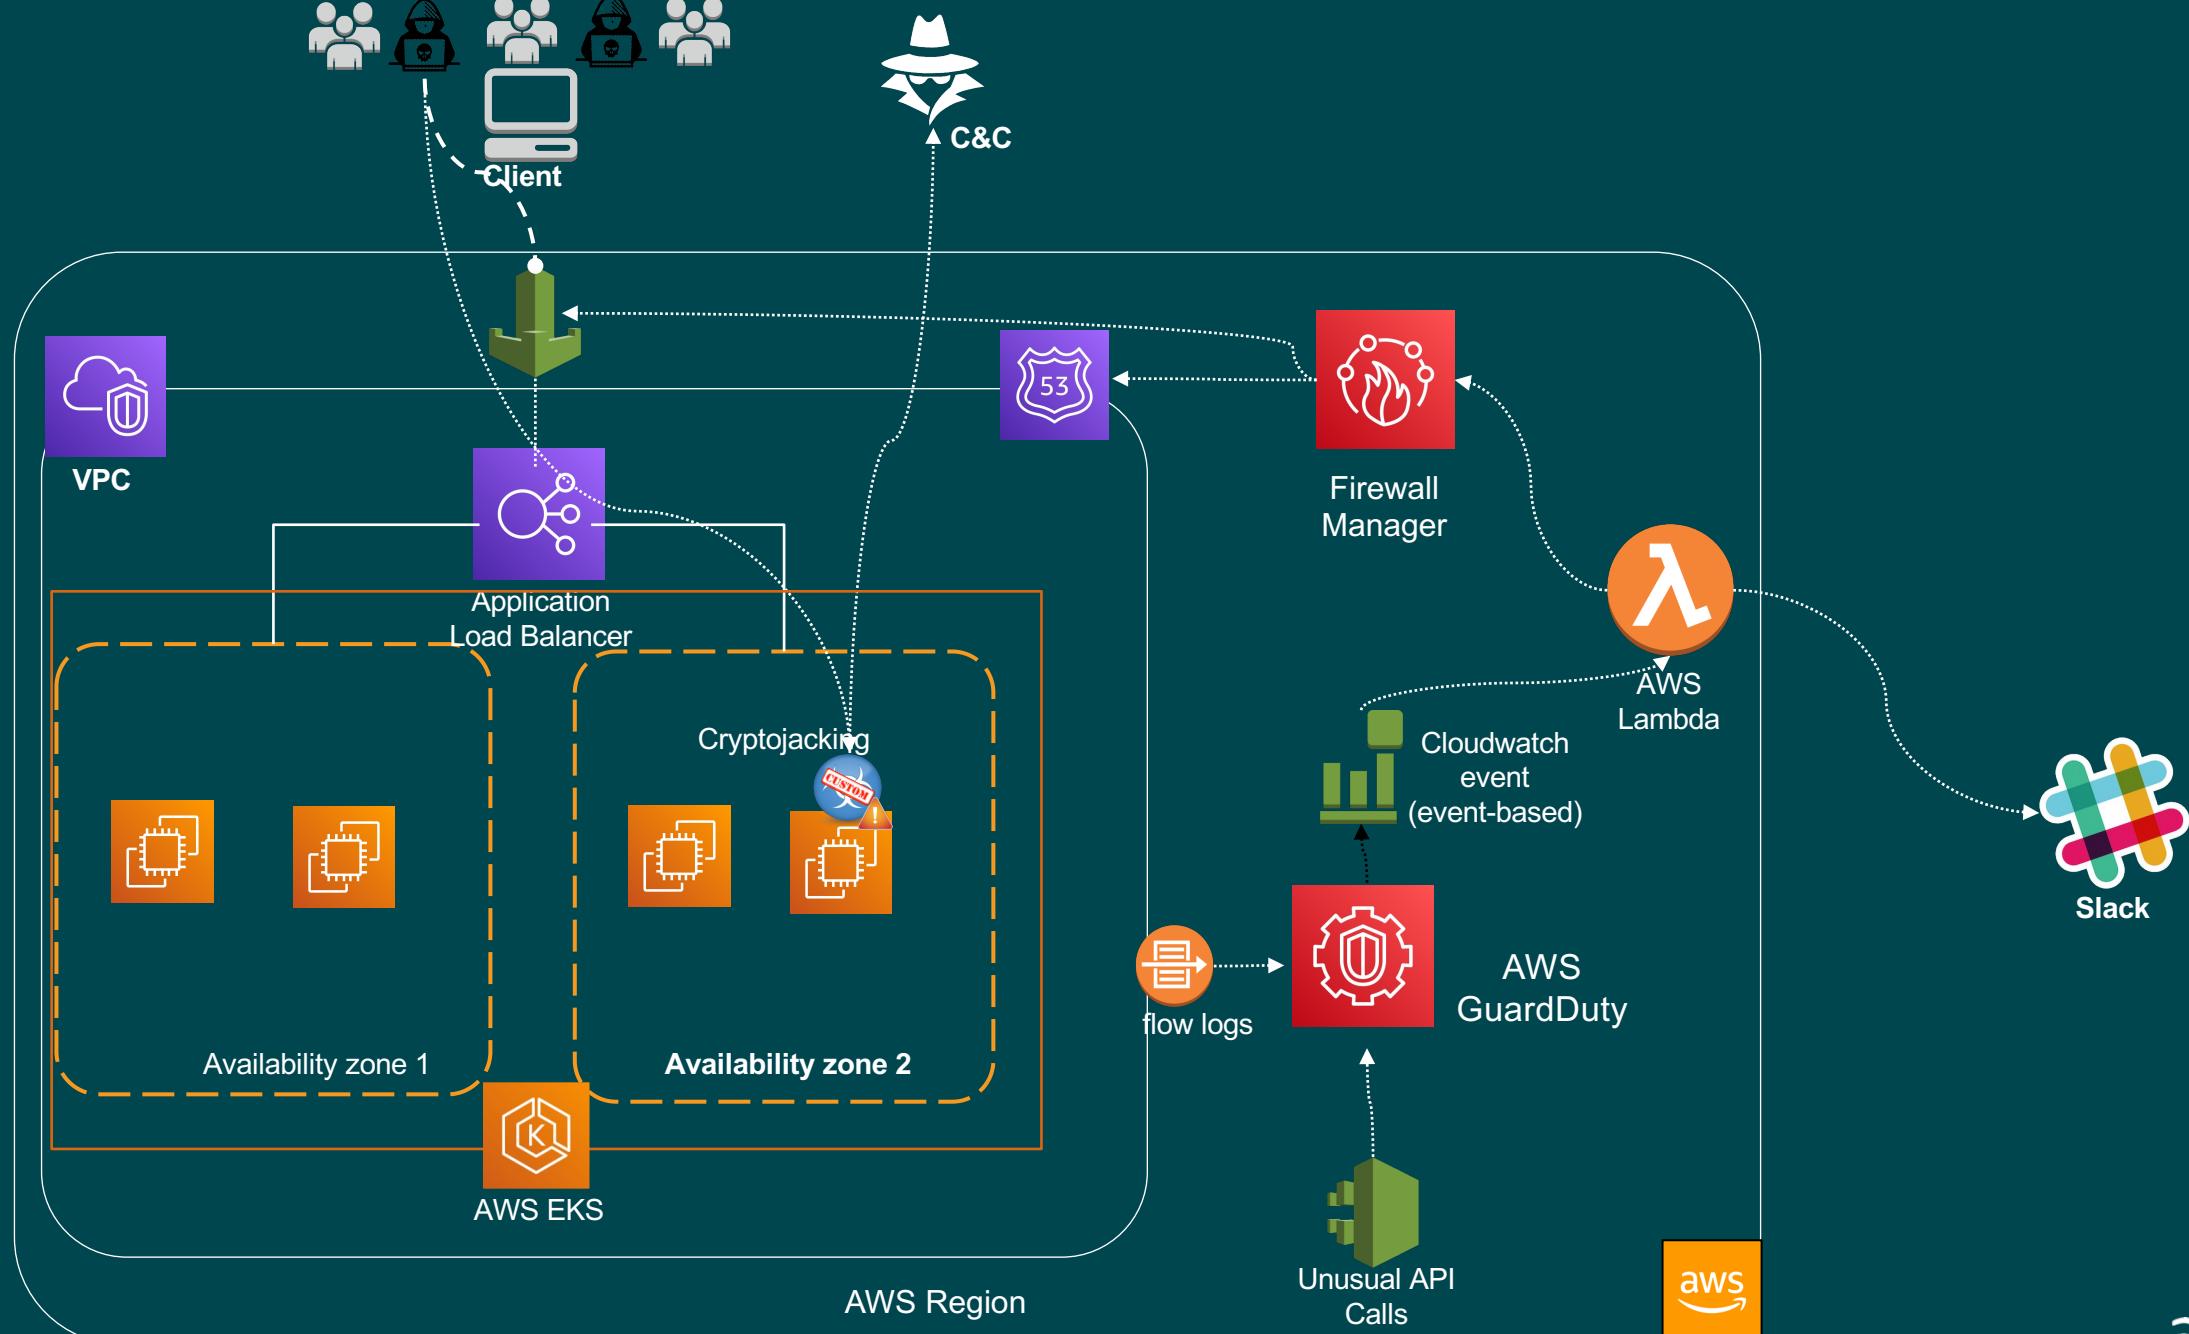


AWS Step Functions
workflows



AWS WAF

and others



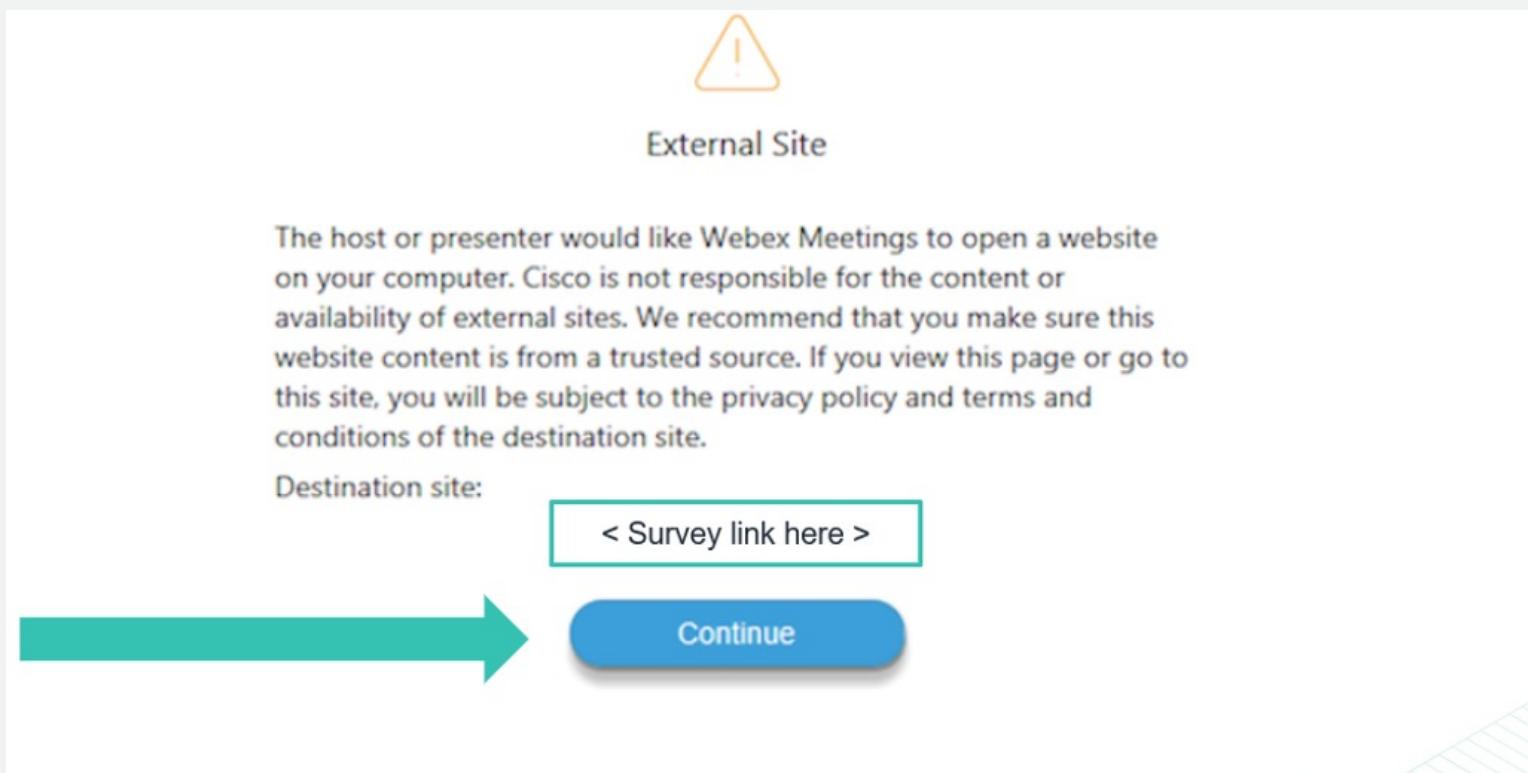
Q&A



© 2023, Amazon Web Services, Inc. or its affiliates.

Exit Survey

We appreciate your feedback! Please complete the survey at the end of this session. You will be automatically redirected via this screen:



Thank you!

Please join us again for another PartnerCast session

<https://aws.amazon.com/partners/training/partnercast/>